

```

import numpy as np

class Softmax(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the Softmax classifier.
        Note that it has shape (C, D) where C is the number of
        classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims) * 0.0001

    def loss(self, X, y):
        """
        Calculates the softmax loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
          that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """
        # Initialize the loss to zero.
        loss = 0.0

        # ===== #
        # YOUR CODE HERE:
        #   Calculate the normalized softmax loss. Store it as the variable loss.
        #   (That is, calculate the sum of the losses of all the training
        #   set margins, and then normalize the loss by the number of
        #   training examples.)
        # ===== #
        pass
        num_classes = self.W.shape[0]
        num_train = X.shape[0]

        losses = np.zeros(num_train)
        for i in np.arange(num_train):
            mar = 0
            for j in np.arange(num_classes):
                mar += np.exp(np.dot(self.W[j], X[i].T))
            margin = np.log(mar) - np.dot(self.W[y[i]], X[i].T)
            losses[i] = margin

        loss = np.sum(losses)/num_train

```

```

# ===== #
# END YOUR CODE HERE
# ===== #

return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
             the gradient of the loss with respect to W.
    """

    # Initialize the loss and gradient to zero.
    loss = 0.0
    grad = np.zeros_like(self.W)

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and the gradient. Store the gradient
    # as the variable grad.
    # ===== #
    pass
    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    # print(self.W.shape) = 10*3073
    # print(X.shape) = 500*3073
    losses = 0.0
    mars = np.zeros(num_train)
    for i in np.arange(num_train):
        mar = 0
        for j in np.arange(num_classes):
            mar += np.exp(np.dot(self.W[j], X[i].T))
        margin = np.log(mar) - np.dot(self.W[y[i]], X[i].T)

        losses += margin

        for j in np.arange(num_classes):
            grad[j] += np.dot(np.exp(np.dot(self.W[j], X[i].T)), X[i].T)/mar
            # print(mar)
            if j == y[i]:
                grad[j] -= X[i].T

    loss = losses/num_train
    grad = grad/num_train
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss, grad

```

```

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
        fxmh = self.loss(X,y) # evaluate f(x - h)
        self.W[ix] = oldval # reset

        grad_numerical = (fxph - fxmh) / (2 * h)
        grad_analytic = your_grad[ix]
        rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) +
            abs(grad_analytic))
        print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical,
            grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero

    # ===== #
    # YOUR CODE HERE:
    # Calculate the softmax loss and gradient WITHOUT any for loops.
    # ===== #
    pass

    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    # print(self.W.shape) = 10*3073
    # print(X.shape) = 500*3073
    losses = 0.0
    mars = np.zeros(num_train)
    scores = np.dot(self.W, X.T)

    mars = np.sum(np.exp(scores), axis = 0)
    margin = np.log(mars) - scores[y, np.arange(0, scores.shape[1])]
    losses = np.sum(margin)
    # print(mars.shape) = (500,)

    # print(np.exp(scores).shape)
    # print(scores.shape)
    # print((np.exp(scores)/mars).shape)

```

```

#     print(X.shape)
grad = np.dot(np.exp(scores)/mars, X)

# eliminate X[i] for j == y[i]
eliminate = np.zeros((num_classes, num_train))
eliminate[y, np.arange(0,scores.shape[1])] = 1
grad -= np.dot(eliminate, X)

loss = losses/num_train
grad = grad/num_train

# ===== #
# END YOUR CODE HERE
# ===== #

return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
        training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
        means that X[i] has label 0 <= c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of
        classes

    self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of
        self.W

    # Run stochastic gradient descent to optimize W
    loss_history = []

    for it in np.arange(num_iters):
        X_batch = None
        y_batch = None

        # ===== #
        # YOUR CODE HERE:
        #     Sample batch_size elements from the training data for use in
        #     gradient descent. After sampling,

```

```

#         - X_batch should have shape: (dim, batch_size)
#         - y_batch should have shape: (batch_size,)
#     The indices should be randomly generated to reduce correlations
#     in the dataset. Use np.random.choice. It's okay to sample with
#     replacement.
# ===== #
idx = np.random.choice(num_train, batch_size)
X_batch = X[idx]
#     print(X_batch.shape)
y_batch = y[idx]
#     print(y_batch.shape)
pass
# ===== #
# END YOUR CODE HERE
# ===== #

# evaluate loss and gradient
loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
loss_history.append(loss)

# ===== #
# YOUR CODE HERE:
#     Update the parameters, self.W, with a gradient step
# ===== #
pass
self.W -= grad*learning_rate
# ===== #
# END YOUR CODE HERE
# ===== #

if verbose and it % 100 == 0:
    print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional
      array of length N, and each element is an integer giving the predicted
      class.
    """
    y_pred = np.zeros(X.shape[1])
    # ===== #
    # YOUR CODE HERE:
    #     Predict the labels given the training data.
    # ===== #
    pass
    scores = np.dot(self.W, X.T)
    y_pred = np.argmax(scores, axis=0)

```

```
# ===== #  
# END YOUR CODE HERE  
# ===== #  
  
return y_pred
```