

```

import numpy as np
import pdb

"""
This code was based off of code from cs231n at Stanford University, and modified for
ece239as at UCLA.
"""

class SVM(object):

    def __init__(self, dims=[10, 3073]):
        self.init_weights(dims=dims)

    def init_weights(self, dims):
        """
        Initializes the weight matrix of the SVM. Note that it has shape (C, D)
        where C is the number of classes and D is the feature size.
        """
        self.W = np.random.normal(size=dims)

    def loss(self, X, y):
        """
        Calculates the SVM loss.

        Inputs have dimension D, there are C classes, and we operate on minibatches
        of N examples.

        Inputs:
        - X: A numpy array of shape (N, D) containing a minibatch of data.
        - y: A numpy array of shape (N,) containing training labels; y[i] = c means
            that X[i] has label c, where 0 <= c < C.

        Returns a tuple of:
        - loss as single float
        """

        # compute the loss and the gradient
        num_classes = self.W.shape[0]
        num_train = X.shape[0]
        loss = 0.0
        hinges = np.zeros(num_train)

        for i in np.arange(num_train):
            # ===== #
            # YOUR CODE HERE:
            # Calculate the normalized SVM loss, and store it as 'loss'.
            # (That is, calculate the sum of the losses of all the training
            # set margins, and then normalize the loss by the number of
            # training examples.)
            # ===== #
            hinge = 0
            for j in np.arange(num_classes):
                if j == y[i]:
                    continue
                hinge_per_j = 1 + np.dot(self.W[j], X[i]) - np.dot(self.W[y[i]], X[i])

```

```

        if hinge_per_j > 0:
            hinge += hinge_per_j
        hinges[i] = hinge

    loss = np.sum(hinges)/num_train

    pass
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    return loss

def loss_and_grad(self, X, y):
    """
    Same as self.loss(X, y), except that it also returns the gradient.

    Output: grad -- a matrix of the same dimensions as W containing
        the gradient of the loss with respect to W.
    """

    # compute the loss and the gradient
    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    # print(X.shape) : 500*3073
    # print(self.W.shape) : 10*3073
    loss = 0.0
    grad = np.zeros_like(self.W)

    hinges = np.zeros(num_train)

    for i in np.arange(num_train):
        # ===== #
        # YOUR CODE HERE:
        # Calculate the SVM loss and the gradient. Store the gradient in
        # the variable grad.
        # ===== #

        hinge = 0
        for j in np.arange(num_classes):
            if j == y[i]:
                continue
            hinge_per_j = 1 + np.dot(self.W[j], X[i]) - np.dot(self.W[y[i]], X[i])

            if hinge_per_j > 0:
                hinge += hinge_per_j
                grad[j] += X[i]
                grad[y[i]] -= X[i]

        hinges[i] = hinge

    loss = np.sum(hinges)

```

```

pass

# ===== #
# END YOUR CODE HERE
# ===== #

loss /= num_train
grad /= num_train

return loss, grad

def grad_check_sparse(self, X, y, your_grad, num_checks=10, h=1e-5):
    """
    sample a few random elements and only return numerical
    in these dimensions.
    """

    for i in np.arange(num_checks):
        ix = tuple([np.random.randint(m) for m in self.W.shape])

        oldval = self.W[ix]
        self.W[ix] = oldval + h # increment by h
        fxph = self.loss(X, y)
        self.W[ix] = oldval - h # decrement by h
        fxmh = self.loss(X,y) # evaluate f(x - h)
        self.W[ix] = oldval # reset

        grad_numerical = (fxph - fxmh) / (2 * h)
        grad_analytic = your_grad[ix]
        rel_error = abs(grad_numerical - grad_analytic) / (abs(grad_numerical) +
            abs(grad_analytic))
        print('numerical: %f analytic: %f, relative error: %e' % (grad_numerical,
            grad_analytic, rel_error))

def fast_loss_and_grad(self, X, y):
    """
    A vectorized implementation of loss_and_grad. It shares the same
    inputs and outputs as loss_and_grad.
    """
    # compute the loss and the gradient
    num_classes = self.W.shape[0]
    num_train = X.shape[0]
    loss = 0.0
    grad = np.zeros(self.W.shape) # initialize the gradient as zero
    # print(X.shape) : 500*3073
    # print(self.W.shape) : 10*3073
    # ===== #
    # YOUR CODE HERE:
    # Calculate the SVM loss WITHOUT any for loops.
    # ===== #
    scores = np.dot(self.W, X.T) # scores.shape = 10*500
    correct_scores = np.ones(scores.shape) * scores[y, np.arange(0,
        scores.shape[1])]
    margin = 1 + scores - correct_scores

```

```

margins = np.maximum(0, margin)
margins[y, np.arange(0, scores.shape[1])] = 0 # y[i] = j position to 0
loss = np.sum(margins)

# ===== #
# END YOUR CODE HERE
# ===== #

# ===== #
# YOUR CODE HERE:
#   Calculate the SVM grad WITHOUT any for loops.
# ===== #

margins_copy = margins
#   print(margins_copy.shape) = 10*500
margins_copy[margins > 0] = 1
margins_copy[margins < 0] = 0
margins_copy[y, np.arange(0, scores.shape[1])] = 0
margins_copy[y, np.arange(0, scores.shape[1])] = -1*np.sum(margins_copy, axis =
    0)
grad = np.dot(margins_copy, X)

# ===== #
# END YOUR CODE HERE
# ===== #
loss /= num_train
grad /= num_train
return loss, grad

def train(self, X, y, learning_rate=1e-3, num_iters=100,
          batch_size=200, verbose=False):
    """
    Train this linear classifier using stochastic gradient descent.

    Inputs:
    - X: A numpy array of shape (N, D) containing training data; there are N
      training samples each of dimension D.
    - y: A numpy array of shape (N,) containing training labels; y[i] = c
      means that X[i] has label 0 ≤ c < C for C classes.
    - learning_rate: (float) learning rate for optimization.
    - num_iters: (integer) number of steps to take when optimizing
    - batch_size: (integer) number of training examples to use at each step.
    - verbose: (boolean) If true, print progress during optimization.

    Outputs:
    A list containing the value of the loss function at each training iteration.
    """
    num_train, dim = X.shape # dim = 3072
    num_classes = np.max(y) + 1 # assume y takes values 0...K-1 where K is number of
      classes

```

```

self.init_weights(dims=[np.max(y) + 1, X.shape[1]]) # initializes the weights of
self.W

# Run stochastic gradient descent to optimize W
loss_history = []

for it in np.arange(num_iters):
    X_batch = None
    y_batch = None

    # ===== #
    # YOUR CODE HERE:
    #   Sample batch_size elements from the training data for use in
    #   gradient descent. After sampling,
    #   - X_batch should have shape: (dim, batch_size)
    #   - y_batch should have shape: (batch_size,)
    #   The indices should be randomly generated to reduce correlations
    #   in the dataset. Use np.random.choice. It's okay to sample with
    #   replacement.
    # ===== #
    idx = np.random.choice(num_train, batch_size)
    X_batch = X[idx]
    y_batch = y[idx]
    # ===== #
    # END YOUR CODE HERE
    # ===== #

    # evaluate loss and gradient
    loss, grad = self.fast_loss_and_grad(X_batch, y_batch)
#     loss, grad = self.loss_and_grad(X_batch, y_batch)
    loss_history.append(loss)

    # ===== #
    # YOUR CODE HERE:
    #   Update the parameters, self.W, with a gradient step
    # ===== #
    self.W -= learning_rate * grad

    # ===== #
    # END YOUR CODE HERE
    # ===== #

    if verbose and it % 100 == 0:
        print('iteration {} / {}: loss {}'.format(it, num_iters, loss))

return loss_history

def predict(self, X):
    """
    Inputs:
    - X: N x D array of training data. Each row is a D-dimensional point.

    Returns:
    - y_pred: Predicted labels for the data in X. y_pred is a 1-dimensional

```

array of length N, and each element is an integer giving the predicted class.

"""

```
y_pred = np.zeros(X.shape[1])
```

```
# ===== #
```

```
# YOUR CODE HERE:
```

```
# Predict the labels given the training data with the parameter self.W.
```

```
# ===== #
```

```
# print(self.W.shape)
```

```
scores = np.dot(self.W, X.T)
```

```
y_pred = np.argmax(scores, axis=0)
```

```
# ===== #
```

```
# END YOUR CODE HERE
```

```
# ===== #
```

```
return y_pred
```