Optimization for Fully Connected Networks

In this notebook, we will implement different optimization rules for gradient descent. We have provided starter code; however, you will need to copy and paste your code from your implementation of the modular fully connected nets in HW #3 to build upon this.

If you did not complete `affine` forward and backwards passes, or `relu` forward and backward passes from HW #3 correctly, you may use another classmate's implementation of these functions for this assignment, or contact us at ece239as.w18@gmail.com.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, and their layer structure. This also includes nndl.fc_net, nndl.layers, and nndl.layer_utils. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

In [1]:

```
## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.fc_net import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_a
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
  """ returns relative error """
  return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

In [2]:

```
# Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
  print('{}: {} '.format(k, data[k].shape))
```

```
X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

Building upon your HW #3 implementation

Copy and paste the following functions from your HW #3 implementation of a modular FC net:

- `affine_forward` in `nndl/layers.py`
- `affine_backward` in `nndl/layers.py`
- `relu_forward` in `nndl/layers.py`
- `relu_backward` in `nndl/layers.py`
- `affine_relu_forward` in `nndl/layer_utils.py`
- `affine_relu_backward` in `nndl/layer_utils.py`
- The `FullyConnectedNet` class in `nndl/fc_net.py`

Test all functions you copy and pasted

In [3]:

```
from nndl.layer_tests import *

affine_forward_test(); print('\n')
affine_backward_test(); print('\n')
relu_forward_test(); print('\n')
relu_backward_test(); print('\n')
affine_relu_test(); print('\n')
fc_net_test()
```

```
If affine_forward function is working, difference should be less than
1e-9:
difference: 9.7698500479884e-10


If affine_backward is working, error should be less than 1e-9::
dx error: 6.196143459966014e-11
dw error: 8.593386693355485e-11
db error: 1.0823119193558136e-10
```

```
If relu_forward function is working, difference should be around 1e-8:
difference: 4.999999798022158e-08


If relu_forward function is working, error should be less than 1e-9:
dx error: 3.2756060069560773e-12


If affine_relu_forward and affine_relu_backward are working, error sho
uld be less than 1e-9::
dx error: 7.750443675078173e-11
dw error: 6.427529014479837e-10
db error: 7.826648993473125e-12


Running check with reg = 0
Initial loss: 2.3032060431964254
W1 relative error: 6.205920037816513e-07
W2 relative error: 9.980549034456402e-06
W3 relative error: 5.011239114978973e-07
b1 relative error: 3.657289365480386e-08
b2 relative error: 6.660245970414922e-09
b3 relative error: 1.3850216166197306e-10
Running check with reg = 3.14
Initial loss: 7.0083127007698565
W1 relative error: 2.1881303675535814e-08
W2 relative error: 5.899353732524129e-08
W3 relative error: 4.0460047202328045e-08
b1 relative error: 1.323365491179507e-07
b2 relative error: 2.2681933406324766e-09
b3 relative error: 2.2474089714334377e-10
```

Training a larger model

In general, proceeding with vanilla stochastic gradient descent to optimize models may be fraught with problems and limitations, as discussed in class. Thus, we implement optimizers that improve on SGD.

SGD + momentum

In the following section, implement SGD with momentum. Read the `nndl/optim.py` API, which is provided by CS231n, and be sure you understand it. After, implement `sgd_momentum` in `nndl/optim.py` . Test your implementation of `sgd_momentum` by running the cell below.

In [4]:

```
from nndl.optim import sgd_momentum

N, D = 4, 5
w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
dw = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)
v = np.linspace(0.6, 0.9, num=N*D).reshape(N, D)

config = {'learning_rate': 1e-3, 'velocity': v}
next_w, _ = sgd_momentum(w, dw, config=config)

expected_next_w = np.asarray([
  [ 0.1406,      0.20738947,  0.27417895,  0.34096842,  0.40775789],
  [ 0.47454737,  0.54133684,  0.60812632,  0.67491579,  0.74170526],
  [ 0.80849474,  0.87528421,  0.94207368,  1.00886316,  1.07565263],
  [ 1.14244211,  1.20923158,  1.27602105,  1.34281053,  1.4096    ]])
expected_velocity = np.asarray([
  [ 0.5406,      0.55475789,  0.56891579, 0.58307368,  0.59723158],
  [ 0.61138947,  0.62554737,  0.63970526,  0.65386316,  0.66802105],
  [ 0.68217895,  0.69633684,  0.71049474,  0.72465263,  0.73881053],
  [ 0.75296842,  0.76712632,  0.78128421,  0.79544211,  0.8096    ]])

print('next_w error: {}'.format(rel_error(next_w, expected_next_w)))
print('velocity error: {}'.format(rel_error(expected_velocity, config['velocity'])))
```

```
next_w error: 8.882347033505819e-09
velocity error: 4.269287743278663e-09
```

SGD + Nesterov momentum

Implement `sgd_nesterov_momentum` in `nndl/optim.py`.

In [5]:

```
from nndl.optim import sgd_nesterov_momentum

N, D = 4, 5
w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
dw = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)
v = np.linspace(0.6, 0.9, num=N*D).reshape(N, D)

config = {'learning_rate': 1e-3, 'velocity': v}
next_w, _ = sgd_nesterov_momentum(w, dw, config=config)

expected_next_w = np.asarray([
  [0.08714,      0.15246105,  0.21778211,  0.28310316,  0.34842421],
  [0.41374526,   0.47906632,  0.54438737,  0.60970842,  0.67502947],
  [0.74035053,   0.80567158,  0.87099263,  0.93631368,  1.00163474],
  [1.06695579,   1.13227684,  1.19759789,  1.26291895,  1.32824   ]])
expected_velocity = np.asarray([
  [ 0.5406,      0.55475789,  0.56891579,  0.58307368,  0.59723158],
  [ 0.61138947,  0.62554737,  0.63970526,  0.65386316,  0.66802105],
  [ 0.68217895,  0.69633684,  0.71049474,  0.72465263,  0.73881053],
  [ 0.75296842,  0.76712632,  0.78128421,  0.79544211,  0.8096    ]])

print('next_w error: {}'.format(rel_error(next_w, expected_next_w)))
print('velocity error: {}'.format(rel_error(expected_velocity, config['velocity'])))
```

```
next_w error: 1.0875186845081027e-08
velocity error: 4.269287743278663e-09
```

Evaluating SGD, SGD+Momentum, and SGD+NesterovMomentum

Run the following cell to train a 6 layer FC net with SGD, SGD+momentum, and SGD+Nesterov momentum. You should see that SGD+momentum achieves a better loss than SGD, and that SGD+Nesterov momentum achieves a slightly better loss (and training accuracy) than SGD+momentum.

In [6]:

```
num_train = 4000
small_data = {
  'X_train': data['X_train'][:num_train],
  'y_train': data['y_train'][:num_train],
  'X_val': data['X_val'],
  'y_val': data['y_val'],
}

solvers = {}

for update_rule in ['sgd', 'sgd_momentum', 'sgd_nesterov_momentum']:
  print('Optimizing with {}'.format(update_rule))
  model = FullyConnectedNet([100, 100, 100, 100, 100], weight_scale=5e-2)
```

```
  solver = Solver(model, small_data,
                  num_epochs=5, batch_size=100,
                  update_rule=update_rule,
                  optim_config={
                    'learning_rate': 1e-2,
                  },
                  verbose=False)
  solvers[update_rule] = solver
  solver.train()
  print

plt.subplot(3, 1, 1)
plt.title('Training loss')
plt.xlabel('Iteration')

plt.subplot(3, 1, 2)
plt.title('Training accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 3)
plt.title('Validation accuracy')
plt.xlabel('Epoch')

for update_rule, solver in solvers.items():
  plt.subplot(3, 1, 1)
  plt.plot(solver.loss_history, 'o', label=update_rule)

  plt.subplot(3, 1, 2)
  plt.plot(solver.train_acc_history, '-o', label=update_rule)

  plt.subplot(3, 1, 3)
  plt.plot(solver.val_acc_history, '-o', label=update_rule)

for i in [1, 2, 3]:
  plt.subplot(3, 1, i)
  plt.legend(loc='upper center', ncol=4)
plt.gcf().set_size_inches(15, 15)
plt.show()
```

```
Optimizing with sgd
Optimizing with sgd_momentum
Optimizing with sgd_nesterov_momentum

/Users/luyutong/Desktop/Yr1Quar2/EE239AS/hw4/hw4_virtualenv/lib/python
3.6/site-packages/matplotlib/cbook/deprecation.py:106: MatplotlibDepre
cationWarning: Adding an axes using the same arguments as a previous a
xes currently reuses the earlier instance.  In a future version, a new
instance will always be created and returned.  Meanwhile, this warning
can be suppressed, and the future behavior ensured, by passing a uniqu
e label to each axes instance.
  warnings.warn(message, mplDeprecation, stacklevel=1)
```

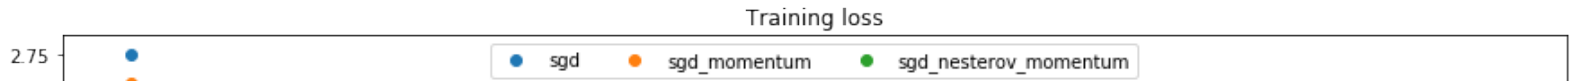

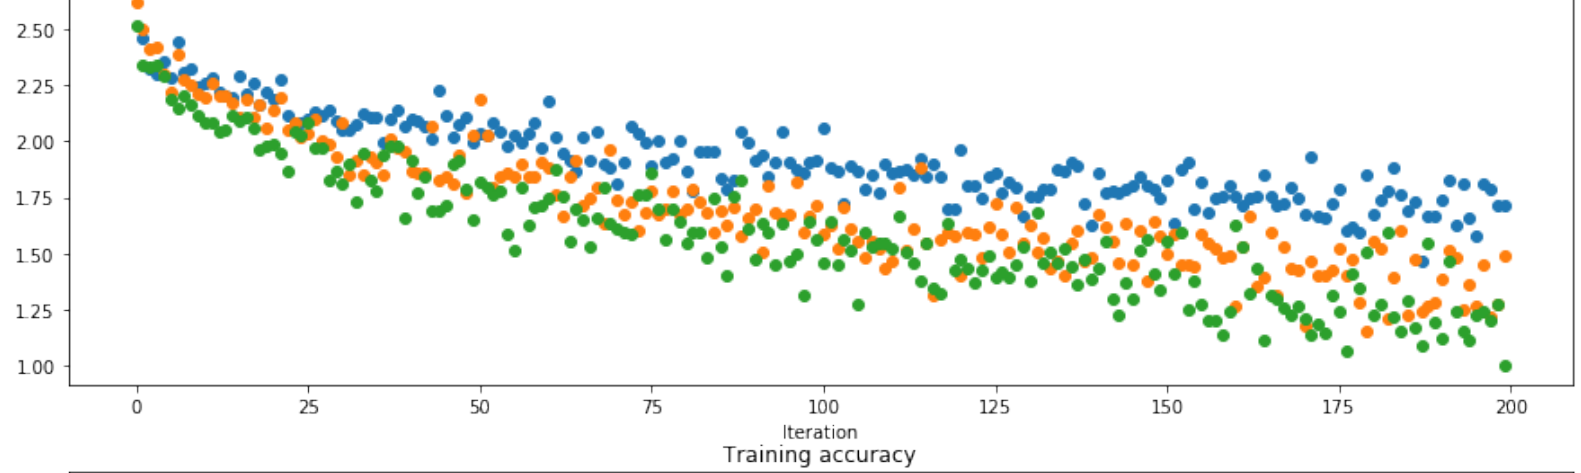

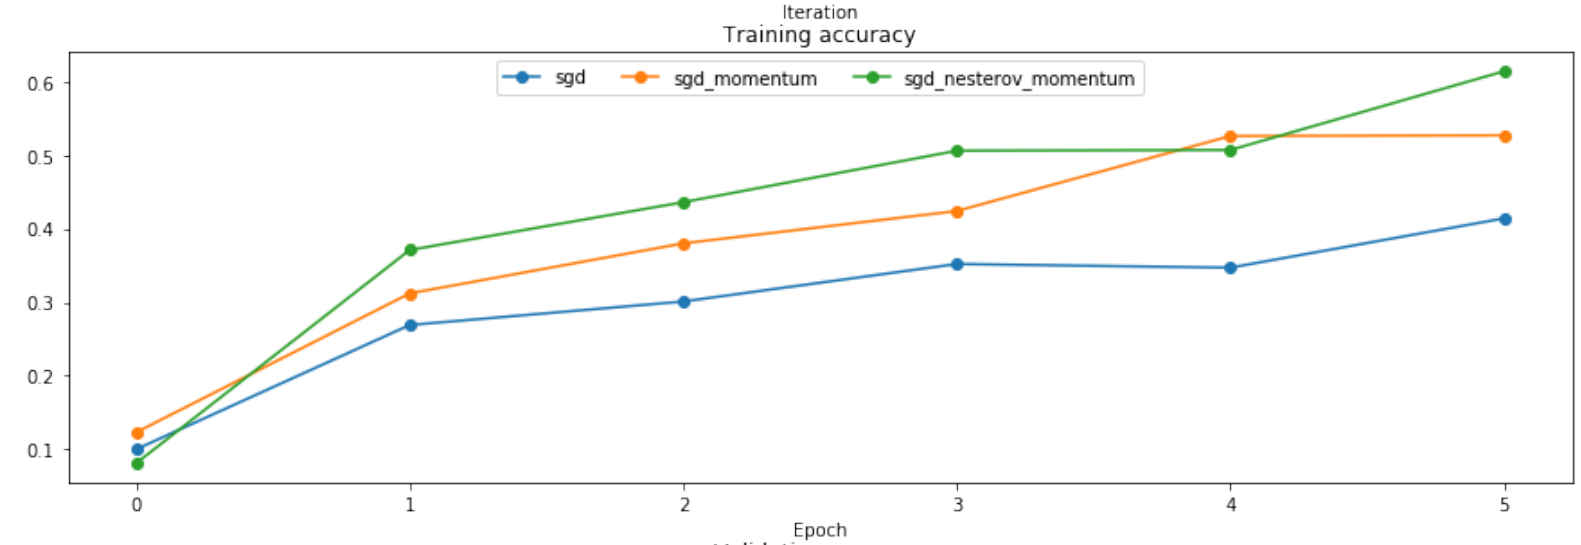

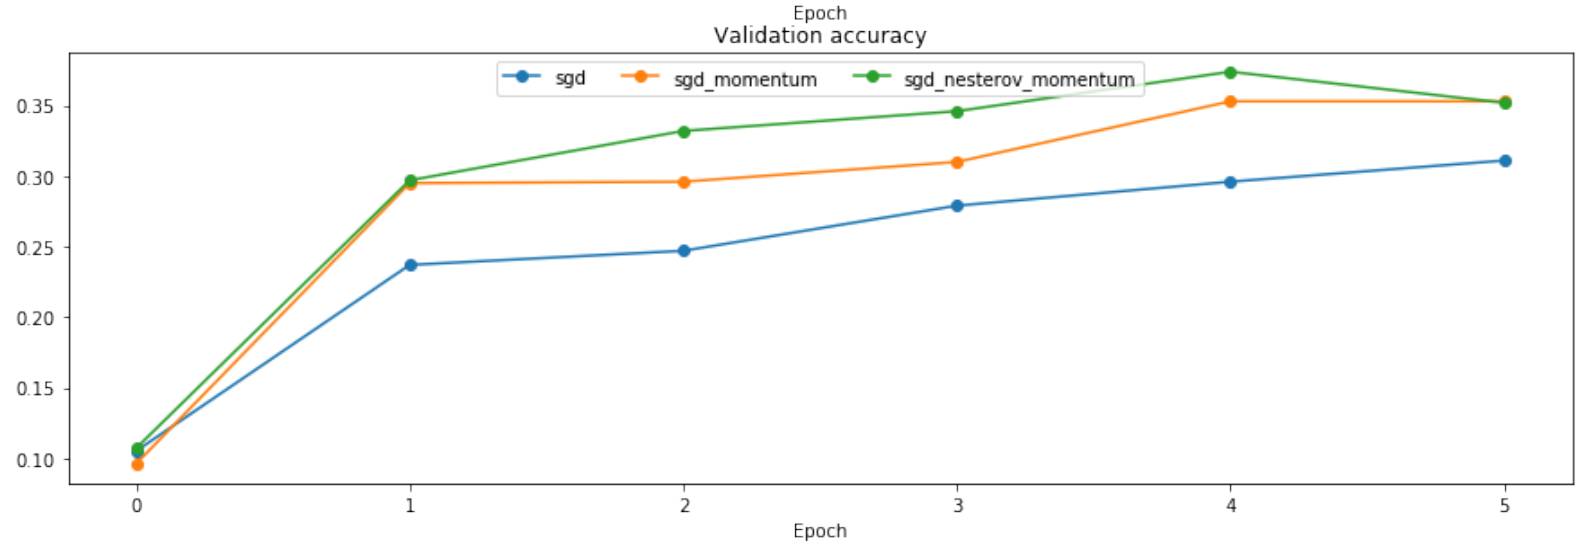

RMSProp

Now we go to techniques that adapt the gradient. Implement `rmsprop` in `nndl/optim.py`. Test your implementation by running the cell below.

In [7]:

```
from nndl.optim import rmsprop

N, D = 4, 5
w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
dw = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)
a = np.linspace(0.6, 0.9, num=N*D).reshape(N, D)

config = {'learning_rate': 1e-2, 'a': a}
next_w, _ = rmsprop(w, dw, config=config)

expected_next_w = np.asarray([
  [-0.39223849, -0.34037513, -0.28849239, -0.23659121, -0.18467247],
  [-0.132737,   -0.08078555, -0.02881884,  0.02316247,  0.07515774],
  [ 0.12716641,  0.17918792,  0.23122175,  0.28326742,  0.33532447],
  [ 0.38739248,  0.43947102,  0.49155973,  0.54365823,  0.59576619]])
expected_cache = np.asarray([
  [ 0.5976,      0.6126277,   0.6277108,   0.64284931,  0.65804321],
  [ 0.67329252,  0.68859723,  0.70395734,  0.71937285,  0.73484377],
  [ 0.75037008,  0.7659518,   0.78158892,  0.79728144,  0.81302936],
  [ 0.82883269,  0.84469141,  0.86060554,  0.87657507,  0.8926    ]])

print('next_w error: {}'.format(rel_error(expected_next_w, next_w)))
print('cache error: {}'.format(rel_error(expected_cache, config['a'])))
```

```
next_w error: 9.524687511038133e-08
cache error: 2.6477955807156126e-09
```

Adaptive moments

Now, implement `adam` in `nndl/optim.py` . Test your implementation by running the cell below.

In [8]:

```
# Test Adam implementation; you should see errors around 1e-7 or less
from nndl.optim import adam

N, D = 4, 5
w = np.linspace(-0.4, 0.6, num=N*D).reshape(N, D)
dw = np.linspace(-0.6, 0.4, num=N*D).reshape(N, D)
v = np.linspace(0.6, 0.9, num=N*D).reshape(N, D)
a = np.linspace(0.7, 0.5, num=N*D).reshape(N, D)

config = {'learning_rate': 1e-2, 'v': v, 'a': a, 't': 5}
next_w, _ = adam(w, dw, config=config)

expected_next_w = np.asarray([
  [-0.40094747, -0.34836187, -0.29577703, -0.24319299, -0.19060977],
  [-0.1380274,  -0.08544591, -0.03286534,  0.01971428,  0.0722929],
  [ 0.1248705,   0.17744702,  0.23002243,  0.28259667,  0.33516969],
  [ 0.38774145,  0.44031188,  0.49288093,  0.54544852,  0.59801459]])
expected_a = np.asarray([
  [ 0.69966,     0.68908382,  0.67851319,  0.66794809,  0.65738853,],
  [ 0.64683452,  0.63628604,  0.6257431,   0.61520571,  0.60467385,],
  [ 0.59414753,  0.58362676,  0.57311152,  0.56260183,  0.55209767,],
  [ 0.54159906,  0.53110598,  0.52061845,  0.51013645,  0.49966,   ]])
expected_v = np.asarray([
  [ 0.48,        0.49947368,  0.51894737,  0.53842105,  0.55789474],
  [ 0.57736842,  0.59684211,  0.61631579,  0.63578947,  0.65526316],
  [ 0.67473684,  0.69421053,  0.71368421,  0.73315789,  0.75263158],
  [ 0.77210526,  0.79157895,  0.81105263,  0.83052632,  0.85      ]])

print('next_w error: {}'.format(rel_error(expected_next_w, next_w)))
print('a error: {}'.format(rel_error(expected_a, config['a'])))
print('v error: {}'.format(rel_error(expected_v, config['v'])))
```

```
next_w error: 1.1395691798535431e-07
a error: 4.208314038113071e-09
v error: 4.214963193114416e-09
```

Comparing SGD, SGD+NesterovMomentum, RMSProp, and Adam

The following code will compare optimization with SGD, Momentum, Nesterov Momentum, RMSProp and Adam. In our code, we find that RMSProp, Adam, and SGD + Nesterov Momentum achieve approximately the same training error after a few training epochs.

In [9]:

```
learning_rates = {'rmsprop': 2e-4, 'adam': 1e-3}

for update_rule in ['adam', 'rmsprop']:
  print('Optimizing with {}'.format(update_rule))
```

```
  model = FullyConnectedNet([100, 100, 100, 100, 100], weight_scale=5e-2)

  solver = Solver(model, small_data,
                  num_epochs=5, batch_size=100,
                  update_rule=update_rule,
                  optim_config={
                    'learning_rate': learning_rates[update_rule]
                  },
                  verbose=False)
  solvers[update_rule] = solver
  solver.train()
  print

plt.subplot(3, 1, 1)
plt.title('Training loss')
plt.xlabel('Iteration')

plt.subplot(3, 1, 2)
plt.title('Training accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 3)
plt.title('Validation accuracy')
plt.xlabel('Epoch')

for update_rule, solver in solvers.items():
  plt.subplot(3, 1, 1)
  plt.plot(solver.loss_history, 'o', label=update_rule)

  plt.subplot(3, 1, 2)
  plt.plot(solver.train_acc_history, '-o', label=update_rule)

  plt.subplot(3, 1, 3)
  plt.plot(solver.val_acc_history, '-o', label=update_rule)

for i in [1, 2, 3]:
  plt.subplot(3, 1, i)
  plt.legend(loc='upper center', ncol=4)
plt.gcf().set_size_inches(15, 15)
plt.show()
```

```
Optimizing with adam
Optimizing with rmsprop

/Users/luyutong/Desktop/Yr1Quar2/EE239AS/hw4/hw4_virtualenv/lib/python
3.6/site-packages/matplotlib/cbook/deprecation.py:106: MatplotlibDepre
cationWarning: Adding an axes using the same arguments as a previous a
xes currently reuses the earlier instance.  In a future version, a new
instance will always be created and returned.  Meanwhile, this warning
can be suppressed, and the future behavior ensured, by passing a uniqu
e label to each axes instance.
  warnings.warn(message, mplDeprecation, stacklevel=1)
```

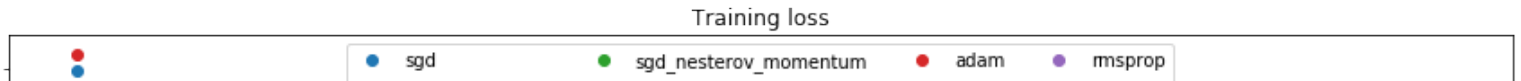

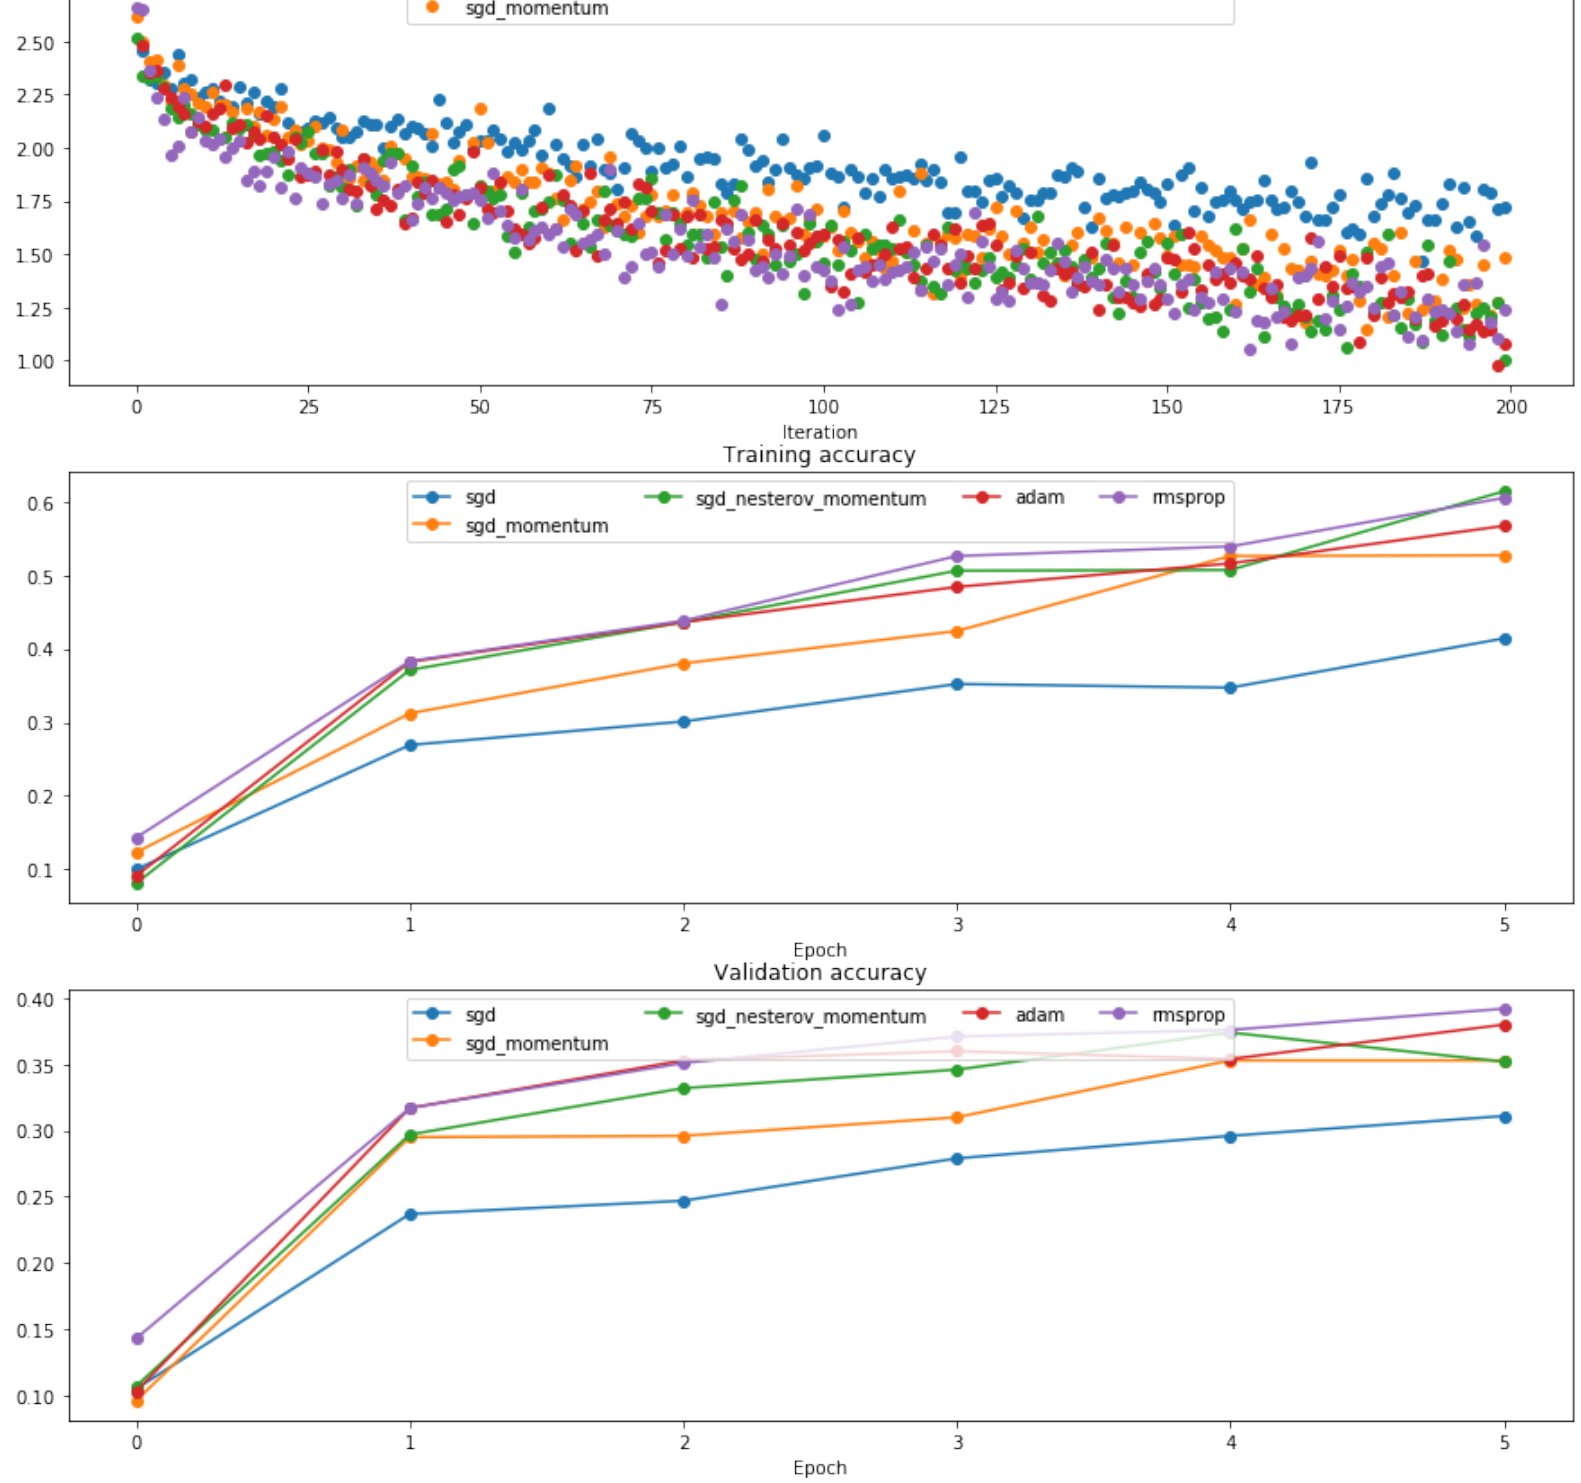

Easier optimization

In the following cell, we'll train a 4 layer neural network having 500 units in each hidden layer with the different optimizers, and find that it is far easier to get up to 50+% performance on CIFAR-10. After we implement batchnorm and dropout, we'll ask you to get 60+% on CIFAR-10.

In [10]:

```
optimizer = 'adam'
best_model = None

layer_dims = [500, 500, 500]
weight_scale = 0.01
learning_rate = 1e-3
lr_decay = 0.9

model = FullyConnectedNet(layer_dims, weight_scale=weight_scale,
                          use_batchnorm=True)
```

```
solver = Solver(model, data,
                num_epochs=10, batch_size=100,
                update_rule=optimizer,
                optim_config={
                  'learning_rate': learning_rate,
                },
                lr_decay=lr_decay,
                verbose=True, print_every=50)
solver.train()
```

```
(Iteration 1 / 4900) loss: 2.287572
(Epoch 0 / 10) train acc: 0.261000; val_acc: 0.218000
(Iteration 51 / 4900) loss: 1.786899
(Iteration 101 / 4900) loss: 1.654803
(Iteration 151 / 4900) loss: 1.477004
(Iteration 201 / 4900) loss: 1.375549
(Iteration 251 / 4900) loss: 1.446279
(Iteration 301 / 4900) loss: 1.572157
(Iteration 351 / 4900) loss: 1.465507
(Iteration 401 / 4900) loss: 1.515942
(Iteration 451 / 4900) loss: 1.523349
(Epoch 1 / 10) train acc: 0.519000; val_acc: 0.475000
(Iteration 501 / 4900) loss: 1.516100
(Iteration 551 / 4900) loss: 1.407207
(Iteration 601 / 4900) loss: 1.236268
(Iteration 651 / 4900) loss: 1.373765
(Iteration 701 / 4900) loss: 1.180377
(Iteration 751 / 4900) loss: 1.299372
(Iteration 801 / 4900) loss: 1.192052
(Iteration 851 / 4900) loss: 1.188752
(Iteration 901 / 4900) loss: 1.218203
(Iteration 951 / 4900) loss: 1.394408
(Epoch 2 / 10) train acc: 0.564000; val_acc: 0.528000
(Iteration 1001 / 4900) loss: 1.370669
(Iteration 1051 / 4900) loss: 1.232192
(Iteration 1101 / 4900) loss: 1.109454
(Iteration 1151 / 4900) loss: 1.079233
(Iteration 1201 / 4900) loss: 1.391881
(Iteration 1251 / 4900) loss: 1.187281
(Iteration 1301 / 4900) loss: 1.287546
(Iteration 1351 / 4900) loss: 1.240658
(Iteration 1401 / 4900) loss: 1.168948
(Iteration 1451 / 4900) loss: 1.199022
(Epoch 3 / 10) train acc: 0.574000; val_acc: 0.525000
(Iteration 1501 / 4900) loss: 0.936441
(Iteration 1551 / 4900) loss: 1.309199
(Iteration 1601 / 4900) loss: 1.175875
(Iteration 1651 / 4900) loss: 0.998241
(Iteration 1701 / 4900) loss: 1.253369
(Iteration 1751 / 4900) loss: 1.213649
(Iteration 1801 / 4900) loss: 1.152420
(Iteration 1851 / 4900) loss: 1.046948
(Iteration 1901 / 4900) loss: 0.988347
```

```
(Iteration 1951 / 4900) loss: 1.086199
(Epoch 4 / 10) train acc: 0.621000; val_acc: 0.539000
(Iteration 2001 / 4900) loss: 1.078610
(Iteration 2051 / 4900) loss: 1.090903
(Iteration 2101 / 4900) loss: 1.088599
(Iteration 2151 / 4900) loss: 0.949241
(Iteration 2201 / 4900) loss: 1.072102
(Iteration 2251 / 4900) loss: 0.770027
(Iteration 2301 / 4900) loss: 0.978615
(Iteration 2351 / 4900) loss: 1.066144
(Iteration 2401 / 4900) loss: 1.022175
(Epoch 5 / 10) train acc: 0.658000; val_acc: 0.530000
(Iteration 2451 / 4900) loss: 0.907661
(Iteration 2501 / 4900) loss: 0.893857
(Iteration 2551 / 4900) loss: 0.799625
(Iteration 2601 / 4900) loss: 1.048767
(Iteration 2651 / 4900) loss: 1.156528
(Iteration 2701 / 4900) loss: 1.160735
(Iteration 2751 / 4900) loss: 0.857120
(Iteration 2801 / 4900) loss: 0.660427
(Iteration 2851 / 4900) loss: 0.844422
(Iteration 2901 / 4900) loss: 1.085005
(Epoch 6 / 10) train acc: 0.703000; val_acc: 0.546000
(Iteration 2951 / 4900) loss: 0.954022
(Iteration 3001 / 4900) loss: 0.843001
(Iteration 3051 / 4900) loss: 0.975890
(Iteration 3101 / 4900) loss: 0.627455
(Iteration 3151 / 4900) loss: 0.956274
(Iteration 3201 / 4900) loss: 0.754262
(Iteration 3251 / 4900) loss: 0.726726
(Iteration 3301 / 4900) loss: 0.749834
(Iteration 3351 / 4900) loss: 0.784334
(Iteration 3401 / 4900) loss: 0.976093
(Epoch 7 / 10) train acc: 0.713000; val_acc: 0.546000
(Iteration 3451 / 4900) loss: 0.678174
(Iteration 3501 / 4900) loss: 0.722890
(Iteration 3551 / 4900) loss: 0.678883
(Iteration 3601 / 4900) loss: 0.664761
(Iteration 3651 / 4900) loss: 0.711431
(Iteration 3701 / 4900) loss: 0.737603
(Iteration 3751 / 4900) loss: 0.816532
(Iteration 3801 / 4900) loss: 0.792706
(Iteration 3851 / 4900) loss: 0.582194
(Iteration 3901 / 4900) loss: 0.748499
(Epoch 8 / 10) train acc: 0.760000; val_acc: 0.545000
(Iteration 3951 / 4900) loss: 0.547727
(Iteration 4001 / 4900) loss: 0.766061
(Iteration 4051 / 4900) loss: 0.761363
(Iteration 4101 / 4900) loss: 0.482564
(Iteration 4151 / 4900) loss: 0.534568
(Iteration 4201 / 4900) loss: 0.687124
(Iteration 4251 / 4900) loss: 0.425396
(Iteration 4301 / 4900) loss: 0.567592
```

```
(Iteration 4351 / 4900) loss: 0.640754
(Iteration 4401 / 4900) loss: 0.551937
(Epoch 9 / 10) train acc: 0.768000; val_acc: 0.562000
(Iteration 4451 / 4900) loss: 0.610165
(Iteration 4501 / 4900) loss: 0.509154
(Iteration 4551 / 4900) loss: 0.565714
(Iteration 4601 / 4900) loss: 0.643105
(Iteration 4651 / 4900) loss: 0.413443
(Iteration 4701 / 4900) loss: 0.476562
(Iteration 4751 / 4900) loss: 0.439850
(Iteration 4801 / 4900) loss: 0.606224
(Iteration 4851 / 4900) loss: 0.514108
(Epoch 10 / 10) train acc: 0.825000; val_acc: 0.562000
```

In [11]:

```
y_test_pred = np.argmax(model.loss(data['X_test']), axis=1)
y_val_pred = np.argmax(model.loss(data['X_val']), axis=1)
print('Validation set accuracy: {}'.format(np.mean(y_val_pred == data['y_val'])))
print('Test set accuracy: {}'.format(np.mean(y_test_pred == data['y_test'])))
```

```
Validation set accuracy: 0.561
Test set accuracy: 0.548
```

In [ ]:

```

```

```
import numpy as np

"""
This code was originally written for CS 231n at Stanford University
(cs231n.stanford.edu).  It has been modified in various areas for use in the
ECE 239AS class at UCLA.  This includes the descriptions of what code to
implement as well as some slight potential changes in variable names to be
consistent with class nomenclature.  We thank Justin Johnson & Serena Yeung for
permission to use this code.  To see the original version, please visit
cs231n.stanford.edu.
"""

"""
This file implements various first-order update rules that are commonly used for
training neural networks. Each update rule accepts current weights and the
gradient of the loss with respect to those weights and produces the next set of
weights. Each update rule has the same interface:

def update(w, dw, config=None):

Inputs:
  - w: A numpy array giving the current weights.
  - dw: A numpy array of the same shape as w giving the gradient of the
    loss with respect to w.
  - config: A dictionary containing hyperparameter values such as learning rate,
    momentum, etc. If the update rule requires caching values over many
    iterations, then config will also hold these cached values.

Returns:
  - next_w: The next point after the update.
  - config: The config dictionary to be passed to the next iteration of the
    update rule.

NOTE: For most update rules, the default learning rate will probably not perform
well; however the default values of the other hyperparameters should work well
for a variety of different problems.

For efficiency, update rules may perform in-place updates, mutating w and
setting next_w equal to w.
"""


def sgd(w, dw, config=None):
  """
  Performs vanilla stochastic gradient descent.

  config format:
  - learning_rate: Scalar learning rate.
  """
  if config is None: config = {}
  config.setdefault('learning_rate', 1e-2)

  w -= config['learning_rate'] * dw
  return w, config
```

```
def sgd_momentum(w, dw, config=None):
  """
  Performs stochastic gradient descent with momentum.

  config format:
  - learning_rate: Scalar learning rate.
  - momentum: Scalar between 0 and 1 giving the momentum value.
    Setting momentum = 0 reduces to sgd.
  - velocity: A numpy array of the same shape as w and dw used to store a moving
    average of the gradients.
  """
  if config is None: config = {}
  config.setdefault('learning_rate', 1e-2)
  config.setdefault('momentum', 0.9) # set momentum to 0.9 if it wasn't there
  v = config.get('velocity', np.zeros_like(w))   # gets velocity, else sets it to
   zero.

  # ================================================================ #
  # YOUR CODE HERE:
  #   Implement the momentum update formula.  Return the updated weights
  #   as next_w, and store the updated velocity as v.
  # ================================================================ #

  pass
  v = config['momentum']*v - config['learning_rate']*dw
  next_w = w + v

  # ================================================================ #
  # END YOUR CODE HERE
  # ================================================================ #

  config['velocity'] = v

  return next_w, config

def sgd_nesterov_momentum(w, dw, config=None):
  """
  Performs stochastic gradient descent with Nesterov momentum.

  config format:
  - learning_rate: Scalar learning rate.
  - momentum: Scalar between 0 and 1 giving the momentum value.
    Setting momentum = 0 reduces to sgd.
  - velocity: A numpy array of the same shape as w and dw used to store a moving
    average of the gradients.
  """
  if config is None: config = {}
  config.setdefault('learning_rate', 1e-2)
  config.setdefault('momentum', 0.9) # set momentum to 0.9 if it wasn't there
  v = config.get('velocity', np.zeros_like(w))   # gets velocity, else sets it to
   zero.
```

```
  # ================================================================ #
  # YOUR CODE HERE:
  #   Implement the momentum update formula.  Return the updated weights
  #   as next_w, and store the updated velocity as v.
  # ================================================================ #
  pass
  v_old = v
  v = config['momentum']*v_old - config['learning_rate']*dw
  next_w = w + v + config['momentum']*(v-v_old)

  # ================================================================ #
  # END YOUR CODE HERE
  # ================================================================ #

  config['velocity'] = v

  return next_w, config

def rmsprop(w, dw, config=None):
  """
  Uses the RMSProp update rule, which uses a moving average of squared gradient
  values to set adaptive per-parameter learning rates.

  config format:
  - learning_rate: Scalar learning rate.
  - decay_rate: Scalar between 0 and 1 giving the decay rate for the squared
    gradient cache.
  - epsilon: Small scalar used for smoothing to avoid dividing by zero.
  - beta: Moving average of second moments of gradients.
  """
  if config is None: config = {}
  config.setdefault('learning_rate', 1e-2)
  config.setdefault('decay_rate', 0.99)
  config.setdefault('epsilon', 1e-8)
  config.setdefault('a', np.zeros_like(w))

  next_w = None

  # ================================================================ #
  # YOUR CODE HERE:
  #   Implement RMSProp.  Store the next value of w as next_w.  You need
  #   to also store in config['a'] the moving average of the second
  #   moment gradients, so they can be used for future gradients. Concretely,
  #   config['a'] corresponds to "a" in the lecture notes.
  # ================================================================ #

  pass
  config['a'] = config['decay_rate']*config['a'] + (1-
   config['decay_rate'])*np.multiply(dw, dw)
  next_w = w - config['learning_rate']*dw/(np.sqrt(config['a'])+config['epsilon'])

  # ================================================================ #
  # END YOUR CODE HERE
  # ================================================================ #
```

```
  return next_w, config


def adam(w, dw, config=None):
  """
  Uses the Adam update rule, which incorporates moving averages of both the
  gradient and its square and a bias correction term.

  config format:
  - learning_rate: Scalar learning rate.
  - beta1: Decay rate for moving average of first moment of gradient.
  - beta2: Decay rate for moving average of second moment of gradient.
  - epsilon: Small scalar used for smoothing to avoid dividing by zero.
  - m: Moving average of gradient.
  - v: Moving average of squared gradient.
  - t: Iteration number.
  """
  if config is None: config = {}
  config.setdefault('learning_rate', 1e-3)
  config.setdefault('beta1', 0.9)
  config.setdefault('beta2', 0.999)
  config.setdefault('epsilon', 1e-8)
  config.setdefault('v', np.zeros_like(w))
  config.setdefault('a', np.zeros_like(w))
  config.setdefault('t', 0)

  next_w = None

  # ================================================================ #
  # YOUR CODE HERE:
  #   Implement Adam.  Store the next value of w as next_w.  You need
  #   to also store in config['a'] the moving average of the second
  #   moment gradients, and in config['v'] the moving average of the
  #   first moments.  Finally, store in config['t'] the increasing time.
  # ================================================================ #

  pass
  config['t'] += 1
  v_new = config['beta1']*config['v'] + (1-config['beta1'])*dw
  a_new = config['beta2']*config['a'] + (1-config['beta2'])*np.multiply(dw, dw)
  v = v_new/(1 - config['beta1']**config['t'])
  a = a_new/(1 - config['beta2']**config['t'])
  next_w = w - config['learning_rate']*v/(np.sqrt(a)+config['epsilon'])
  config['v'] = v_new
  config['a'] = a_new
  # ================================================================ #
  # END YOUR CODE HERE
  # ================================================================ #

  return next_w, config
```

Batch Normalization

In this notebook, you will implement the batch normalization layers of a neural network to increase its performance. If you have any confusion, please review the details of batch normalization from the lecture notes.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, and their layer structure. This also includes nndl.fc_net, nndl.layers, and nndl.layer_utils. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

In [1]:

```
## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.fc_net import *
from nndl.layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_a
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
  """ returns relative error """
  return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

In [2]:

```
# Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
  print('{}: {} '.format(k, data[k].shape))
```

```
X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

Batchnorm forward pass

Implement the training time batchnorm forward pass, `batchnorm_forward` , in `nndl/layers.py` . After that, test your implementation by running the following cell.

In [3]:

```
# Check the training-time forward pass by checking means and variances
# of features both before and after batch normalization

# Simulate the forward pass for a two-layer network
N, D1, D2, D3 = 200, 50, 60, 3
X = np.random.randn(N, D1)
W1 = np.random.randn(D1, D2)
W2 = np.random.randn(D2, D3)
a = np.maximum(0, X.dot(W1)).dot(W2)

print('Before batch normalization:')
print('  means: ', a.mean(axis=0))
print('  stds: ', a.std(axis=0))

# Means should be close to zero and stds close to one
print('After batch normalization (gamma=1, beta=0)')
a_norm, _ = batchnorm_forward(a, np.ones(D3), np.zeros(D3), {'mode': 'train'})
print('  mean: ', a_norm.mean(axis=0))
print('  std: ', a_norm.std(axis=0))

# Now means should be close to beta and stds close to gamma
gamma = np.asarray([1.0, 2.0, 3.0])
beta = np.asarray([11.0, 12.0, 13.0])
a_norm, _ = batchnorm_forward(a, gamma, beta, {'mode': 'train'})
print('After batch normalization (nontrivial gamma, beta)')
print('  means: ', a_norm.mean(axis=0))
print('  stds: ', a_norm.std(axis=0))
```

```
Before batch normalization:
  means:  [-34.63913235   8.02031228  27.81437078]
  stds:  [36.0004351  34.24710459 44.10177699]
After batch normalization (gamma=1, beta=0)
  mean:  [5.02931030e-16 2.52575738e-17 1.03316861e-16]
  std:  [1. 1. 1.]
After batch normalization (nontrivial gamma, beta)
  means:  [11. 12. 13.]
  stds:  [1.         1.99999999 2.99999999]
```

Implement the testing time batchnorm forward pass, `batchnorm_forward` , in `nndl/layers.py` . After that, test your implementation by running the following cell.

In [4]:

```
# Check the test-time forward pass by running the training-time
# forward pass many times to warm up the running averages, and then
# checking the means and variances of activations after a test-time
# forward pass.

N, D1, D2, D3 = 200, 50, 60, 3
W1 = np.random.randn(D1, D2)
W2 = np.random.randn(D2, D3)

bn_param = {'mode': 'train'}
gamma = np.ones(D3)
beta = np.zeros(D3)
for t in np.arange(50):
  X = np.random.randn(N, D1)
  a = np.maximum(0, X.dot(W1)).dot(W2)
  batchnorm_forward(a, gamma, beta, bn_param)
bn_param['mode'] = 'test'
X = np.random.randn(N, D1)
a = np.maximum(0, X.dot(W1)).dot(W2)
a_norm, _ = batchnorm_forward(a, gamma, beta, bn_param)

# Means should be close to zero and stds close to one, but will be
# noisier than training-time forward passes.
print('After batch normalization (test-time):')
print('  means: ', a_norm.mean(axis=0))
print('  stds: ', a_norm.std(axis=0))
```

```
After batch normalization (test-time):
  means:  [ 0.07872688 -0.04093016 -0.08922362]
  stds:  [1.0576903  0.99762145 0.95151787]
```

Batchnorm backward pass

Implement the backward pass for the batchnorm layer, `batchnorm_backward` in `nndl/layers.py`. Check your implementation by running the following cell.

In [5]:

```
# Gradient check batchnorm backward pass

N, D = 4, 5
x = 5 * np.random.randn(N, D) + 12
gamma = np.random.randn(D)
beta = np.random.randn(D)
dout = np.random.randn(N, D)

bn_param = {'mode': 'train'}
fx = lambda x: batchnorm_forward(x, gamma, beta, bn_param)[0]
fg = lambda a: batchnorm_forward(x, gamma, beta, bn_param)[0]
fb = lambda b: batchnorm_forward(x, gamma, beta, bn_param)[0]

dx_num = eval_numerical_gradient_array(fx, x, dout)
da_num = eval_numerical_gradient_array(fg, gamma, dout)
db_num = eval_numerical_gradient_array(fb, beta, dout)

_, cache = batchnorm_forward(x, gamma, beta, bn_param)
dx, dgamma, dbeta = batchnorm_backward(dout, cache)
print('dx error: ', rel_error(dx_num, dx))
print('dgamma error: ', rel_error(da_num, dgamma))
print('dbeta error: ', rel_error(db_num, dbeta))
```

```
dx error:  5.438826357167099e-10
dgamma error:  2.9894815332189l9e-12
dbeta error:  3.275645486029122e-12
```

Implement a fully connected neural network with batchnorm layers

Modify the `FullyConnectedNet()` class in `nndl/fc_net.py` to incorporate batchnorm layers. You will need to modify the class in the following areas:

(1) The gammas and betas need to be initialized to 1's and 0's respectively in `__init__` .

(2) The `batchnorm_forward` layer needs to be inserted between each affine and relu layer (except in the output layer) in a forward pass computation in `loss` . You may find it helpful to write an `affine_batchnorm_relu()` layer in `nndl/layer_utils.py` although this is not necessary.

(3) The `batchnorm_backward` layer has to be appropriately inserted when calculating gradients.

After you have done the appropriate modifications, check your implementation by running the following cell.

Note, while the relative error for W3 should be small, as we backprop gradients more, you may find the relative error increases. Our relative error for W1 is on the order of 1e-4.

In [6]:

```
N, D, H1, H2, C = 2, 15, 20, 30, 10
X = np.random.randn(N, D)
y = np.random.randint(C, size=(N,))

for reg in [0, 3.14]:
  print('Running check with reg = ', reg)
  model = FullyConnectedNet([H1, H2], input_dim=D, num_classes=C,
                            reg=reg, weight_scale=5e-2, dtype=np.float64,
                            use_batchnorm=True)

  loss, grads = model.loss(X, y)
  print('Initial loss: ', loss)

  for name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    grad_num = eval_numerical_gradient(f, model.params[name], verbose=False, h=1e-5)
    print('{} relative error: {}'.format(name, rel_error(grad_num, grads[name])))
  if reg == 0: print('\n')
```

```
Running check with reg =  0
Initial loss:  2.2211738368223894
W1 relative error: 5.925251839604438e-06
W2 relative error: 0.0002780564470580 82
W3 relative error: 4.745842505044503e-07
b1 relative error: 0.004440536827132745
b2 relative error: 0.0022204460492503126
b3 relative error: 4.779071408552745e-07
beta1 relative error: 4.73619408597463e-07
beta2 relative error: 4.76028832939637e-07
gamma1 relative error: 4.7182194817663683e-07
gamma2 relative error: 4.788565027705718e-07


Running check with reg =  3.14
Initial loss:  6.692697238391027
W1 relative error: 7.697131670776335e-05
W2 relative error: 6.808932503531371e-06
W3 relative error: 4.102207598446421e-05
b1 relative error: 1.7763568394002505e-07
b2 relative error: 2.220446049250313e-08
b3 relative error: 4.7210835730734625e-07
beta1 relative error: 6.74498069818721e-07
beta2 relative error: 4.5888452928750607e-07
gamma1 relative error: 6.76241015540894e-07
gamma2 relative error: 4.539803822469235e-07
```

Training a deep fully connected network with batch normalization.

To see if batchnorm helps, let's train a deep neural network with and without batch normalization.

In [7]:

```
# Try training a very deep net with batchnorm
hidden_dims = [100, 100, 100, 100, 100]

num_train = 1000
small_data = {
  'X_train': data['X_train'][:num_train],
  'y_train': data['y_train'][:num_train],
  'X_val': data['X_val'],
  'y_val': data['y_val'],
}

weight_scale = 2e-2
bn_model = FullyConnectedNet(hidden_dims, weight_scale=weight_scale, use_batchnorm=T
model = FullyConnectedNet(hidden_dims, weight_scale=weight_scale, use_batchnorm=Fals

bn_solver = Solver(bn_model, small_data,
                num_epochs=10, batch_size=50,
                update_rule='adam',
                optim_config={
                  'learning_rate': 1e-3,
                },
                verbose=True, print_every=200)
bn_solver.train()

solver = Solver(model, small_data,
                num_epochs=10, batch_size=50,
                update_rule='adam',
                optim_config={
                  'learning_rate': 1e-3,
                },
                verbose=True, print_every=200)
solver.train()
```

```
(Iteration 1 / 200) loss: 2.305632
(Epoch 0 / 10) train acc: 0.142000; val_acc: 0.135000
(Epoch 1 / 10) train acc: 0.346000; val_acc: 0.269000
(Epoch 2 / 10) train acc: 0.419000; val_acc: 0.296000
(Epoch 3 / 10) train acc: 0.514000; val_acc: 0.314000
(Epoch 4 / 10) train acc: 0.553000; val_acc: 0.309000
(Epoch 5 / 10) train acc: 0.596000; val_acc: 0.333000
(Epoch 6 / 10) train acc: 0.631000; val_acc: 0.300000
(Epoch 7 / 10) train acc: 0.704000; val_acc: 0.337000
(Epoch 8 / 10) train acc: 0.756000; val_acc: 0.322000
(Epoch 9 / 10) train acc: 0.775000; val_acc: 0.321000
```

```
(Epoch 10 / 10) train acc: 0.805000; val_acc: 0.326000
(Iteration 1 / 200) loss: 2.302068
(Epoch 0 / 10) train acc: 0.122000; val_acc: 0.122000
(Epoch 1 / 10) train acc: 0.238000; val_acc: 0.211000
(Epoch 2 / 10) train acc: 0.229000; val_acc: 0.175000
(Epoch 3 / 10) train acc: 0.332000; val_acc: 0.264000
(Epoch 4 / 10) train acc: 0.335000; val_acc: 0.233000
(Epoch 5 / 10) train acc: 0.409000; val_acc: 0.278000
(Epoch 6 / 10) train acc: 0.471000; val_acc: 0.285000
(Epoch 7 / 10) train acc: 0.514000; val_acc: 0.316000
(Epoch 8 / 10) train acc: 0.504000; val_acc: 0.289000
(Epoch 9 / 10) train acc: 0.579000; val_acc: 0.298000
(Epoch 10 / 10) train acc: 0.614000; val_acc: 0.305000
```

In [8]:

```
plt.subplot(3, 1, 1)
plt.title('Training loss')
plt.xlabel('Iteration')

plt.subplot(3, 1, 2)
plt.title('Training accuracy')
plt.xlabel('Epoch')

plt.subplot(3, 1, 3)
plt.title('Validation accuracy')
plt.xlabel('Epoch')


plt.subplot(3, 1, 1)
plt.plot(solver.loss_history, 'o', label='baseline')
plt.plot(bn_solver.loss_history, 'o', label='batchnorm')

plt.subplot(3, 1, 2)
plt.plot(solver.train_acc_history, '-o', label='baseline')
plt.plot(bn_solver.train_acc_history, '-o', label='batchnorm')

plt.subplot(3, 1, 3)
plt.plot(solver.val_acc_history, '-o', label='baseline')
plt.plot(bn_solver.val_acc_history, '-o', label='batchnorm')

for i in [1, 2, 3]:
  plt.subplot(3, 1, i)
  plt.legend(loc='upper center', ncol=4)
plt.gcf().set_size_inches(15, 15)
plt.show()
```

```
/Users/luyutong/Desktop/Yr1Quar2/EE239AS/hw4/hw4_virtualenv/lib/python
3.6/site-packages/matplotlib/cbook/deprecation.py:106: MatplotlibDepre
cationWarning: Adding an axes using the same arguments as a previous a
xes currently reuses the earlier instance.  In a future version, a new
instance will always be created and returned.  Meanwhile, this warning
can be suppressed, and the future behavior ensured, by passing a uniqu
e label to each axes instance.
```

```
warnings.warn(message, mplDeprecation, stacklevel=1)
```

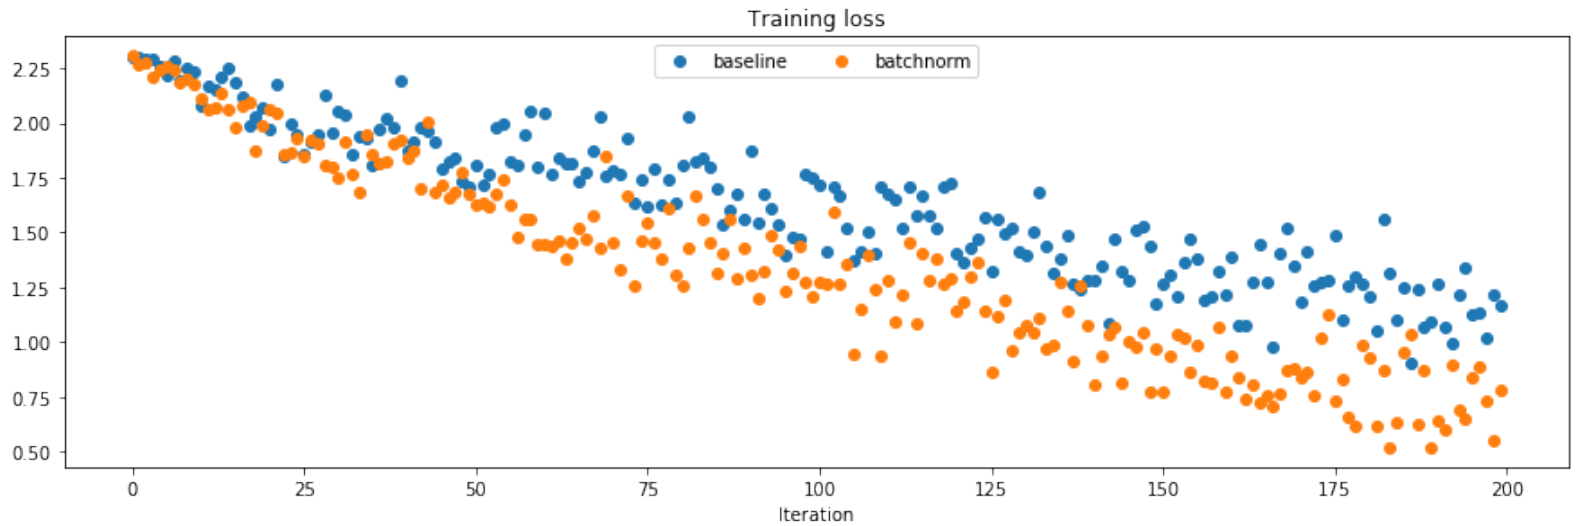

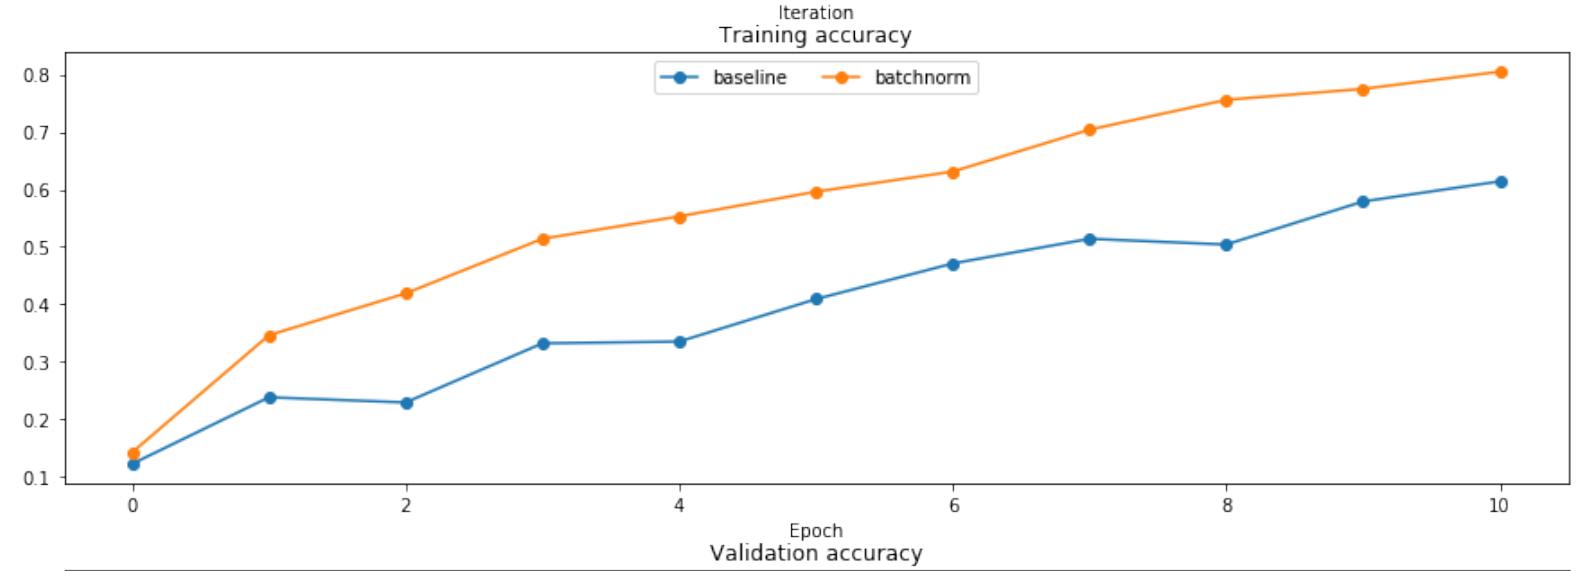

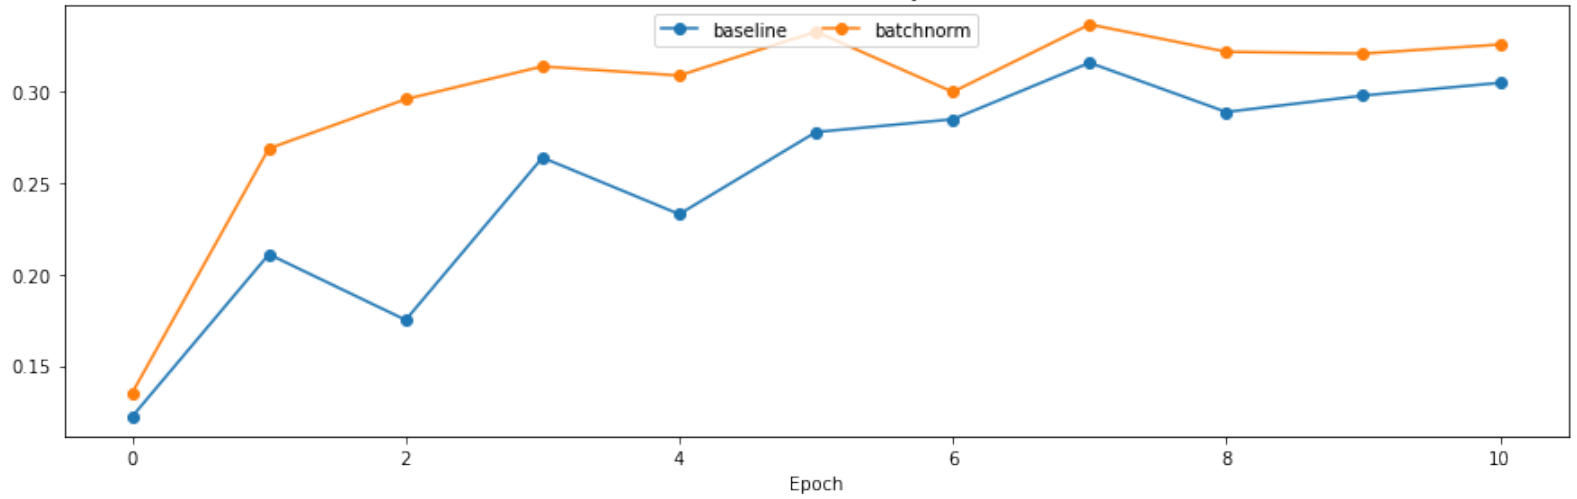

Batchnorm and initialization

The following cells run an experiment where for a deep network, the initialization is varied. We do training for when batchnorm layers are and are not included.

In [9]:

```
# Try training a very deep net with batchnorm
hidden_dims = [50, 50, 50, 50, 50, 50, 50]

num_train = 1000
small_data = {
  'X_train': data['X_train'][:num_train],
  'y_train': data['y_train'][:num_train],
```

```
  'x_val': data['x_val'],
  'y_val': data['y_val'],
}

bn_solvers = {}
solvers = {}
weight_scales = np.logspace(-4, 0, num=20)
for i, weight_scale in enumerate(weight_scales):
  print('Running weight scale {} / {}'.format(i + 1, len(weight_scales)))
  bn_model = FullyConnectedNet(hidden_dims, weight_scale=weight_scale, use_batchnorm
  model = FullyConnectedNet(hidden_dims, weight_scale=weight_scale, use_batchnorm=Fa

  bn_solver = Solver(bn_model, small_data,
                  num_epochs=10, batch_size=50,
                  update_rule='adam',
                  optim_config={
                    'learning_rate': 1e-3,
                  },
                  verbose=False, print_every=200)
  bn_solver.train()
  bn_solvers[weight_scale] = bn_solver

  solver = Solver(model, small_data,
                  num_epochs=10, batch_size=50,
                  update_rule='adam',
                  optim_config={
                    'learning_rate': 1e-3,
                  },
                  verbose=False, print_every=200)
  solver.train()
  solvers[weight_scale] = solver
```

```
Running weight scale 1 / 20
Running weight scale 2 / 20
Running weight scale 3 / 20
Running weight scale 4 / 20
Running weight scale 5 / 20
Running weight scale 6 / 20
Running weight scale 7 / 20
Running weight scale 8 / 20
Running weight scale 9 / 20
Running weight scale 10 / 20
Running weight scale 11 / 20
Running weight scale 12 / 20
Running weight scale 13 / 20
Running weight scale 14 / 20
Running weight scale 15 / 20
Running weight scale 16 / 20
Running weight scale 17 / 20
Running weight scale 18 / 20
Running weight scale 19 / 20
Running weight scale 20 / 20
```

In [10]:

```
# Plot results of weight scale experiment
best_train_accs, bn_best_train_accs = [], []
best_val_accs, bn_best_val_accs = [], []
final_train_loss, bn_final_train_loss = [], []

for ws in weight_scales:
  best_train_accs.append(max(solvers[ws].train_acc_history))
  bn_best_train_accs.append(max(bn_solvers[ws].train_acc_history))

  best_val_accs.append(max(solvers[ws].val_acc_history))
  bn_best_val_accs.append(max(bn_solvers[ws].val_acc_history))

  final_train_loss.append(np.mean(solvers[ws].loss_history[-100:]))
  bn_final_train_loss.append(np.mean(bn_solvers[ws].loss_history[-100:]))

plt.subplot(3, 1, 1)
plt.title('Best val accuracy vs weight initialization scale')
plt.xlabel('Weight initialization scale')
plt.ylabel('Best val accuracy')
plt.semilogx(weight_scales, best_val_accs, '-o', label='baseline')
plt.semilogx(weight_scales, bn_best_val_accs, '-o', label='batchnorm')
plt.legend(ncol=2, loc='lower right')

plt.subplot(3, 1, 2)
plt.title('Best train accuracy vs weight initialization scale')
plt.xlabel('Weight initialization scale')
plt.ylabel('Best training accuracy')
plt.semilogx(weight_scales, best_train_accs, '-o', label='baseline')
plt.semilogx(weight_scales, bn_best_train_accs, '-o', label='batchnorm')
plt.legend()

plt.subplot(3, 1, 3)
plt.title('Final training loss vs weight initialization scale')
plt.xlabel('Weight initialization scale')
plt.ylabel('Final training loss')
plt.semilogx(weight_scales, final_train_loss, '-o', label='baseline')
plt.semilogx(weight_scales, bn_final_train_loss, '-o', label='batchnorm')
plt.legend()

plt.gcf().set_size_inches(10, 15)
plt.show()
```

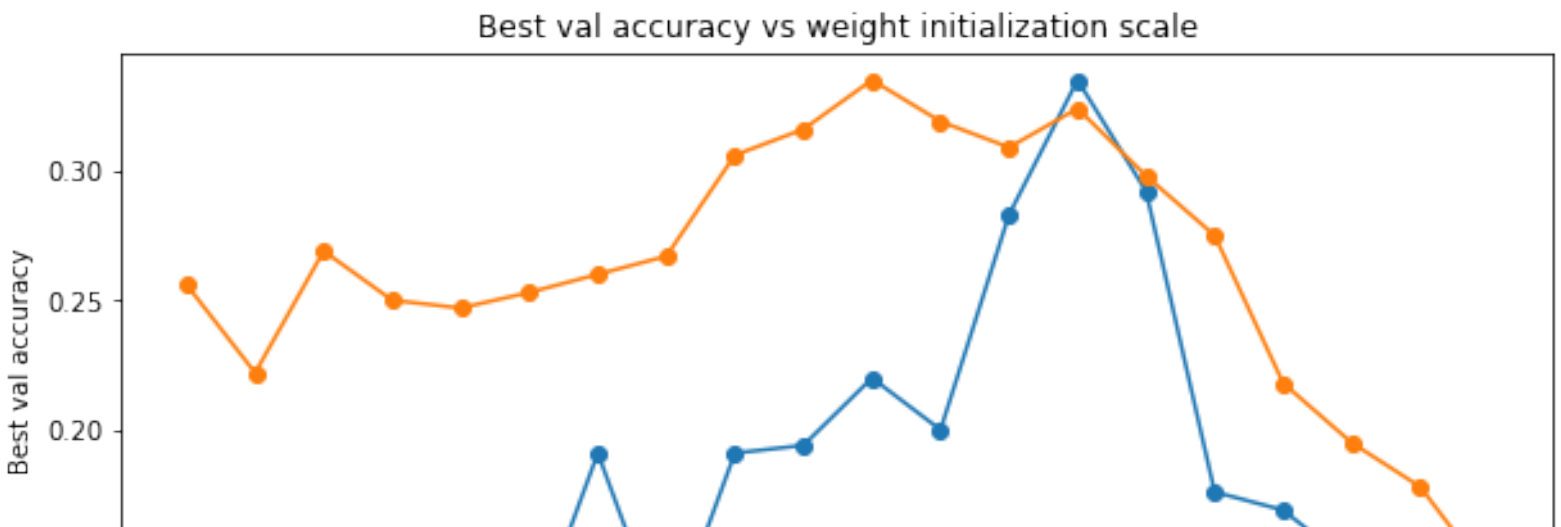

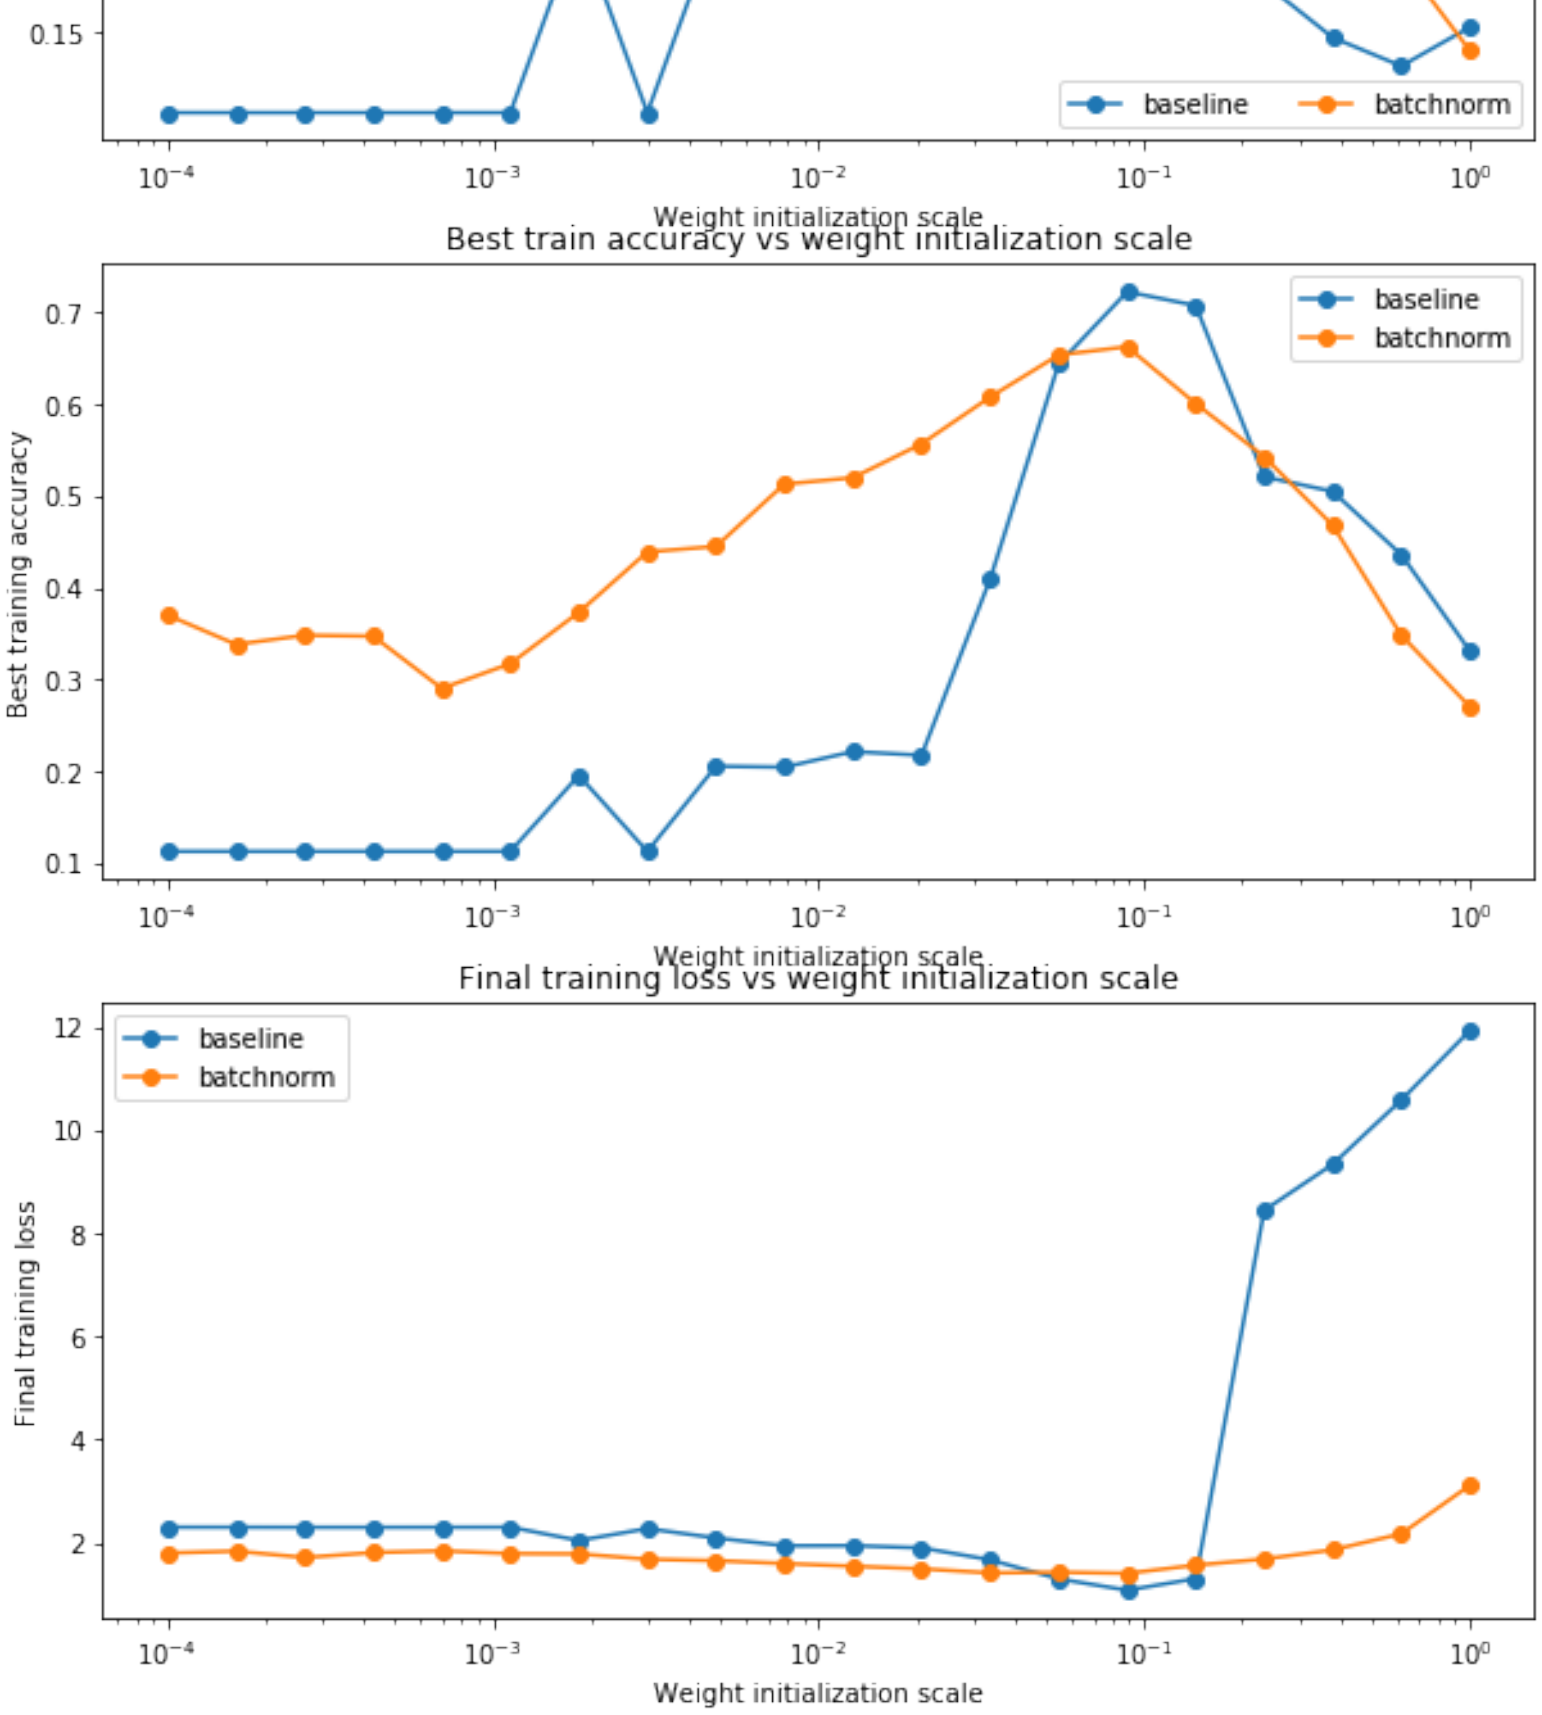

Question:

In the cell below, summarize the findings of this experiment, and WHY these results make sense.

Answer:

It can be observed that batchnorm gives better performance than the baseline overall. When the weight initialization scale is pretty small(~1e-4) or large(~1), both training accuracy and validation accuracy for the baseline are ver low and have high training losses. After applying batchnorm, both the training and validation accuracies will be improved and the final training loss will be reduced. Therefore, it can be concluded that batch normalization is robust and weight initialization scale should be neither too large or too small, otherwise it will impede the training and validation accuracies.

The cause of the above observation is that when the weight initialization scale is small, it is easy to kill the activation in the latter hidden layers, then the output for every hidden layer may be the same. Similarly, when the weight initialization scale is large, the network may be saturating too fast. Therefore, only choosing appropriate weight initialization scale will ensure the process of training. As for the robustness for the batch normalization, since after every activation of the hidden layer, it will transfer the output to be a distribution with zero mean and variance being one. Therefore, it will make the network not that symmetrical and make it not die or saturate too fast.

In [ ]:

```python
import numpy as np
import pdb

"""
This code was originally written for CS 231n at Stanford University
(cs231n.stanford.edu).  It has been modified in various areas for use in the
ECE 239AS class at UCLA.  This includes the descriptions of what code to
implement as well as some slight potential changes in variable names to be
consistent with class nomenclature.  We thank Justin Johnson & Serena Yeung for
permission to use this code.  To see the original version, please visit
cs231n.stanford.edu.
"""

def affine_forward(x, w, b):
  """
  Computes the forward pass for an affine (fully-connected) layer.

  The input x has shape (N, d_1, ..., d_k) and contains a minibatch of N
  examples, where each example x[i] has shape (d_1, ..., d_k). We will
  reshape each input into a vector of dimension D = d_1 * ... * d_k, and
  then transform it to an output vector of dimension M.

  Inputs:
  - x: A numpy array containing input data, of shape (N, d_1, ..., d_k)
  - w: A numpy array of weights, of shape (D, M)
  - b: A numpy array of biases, of shape (M,)

  Returns a tuple of:
  - out: output, of shape (N, M)
  - cache: (x, w, b)
  """

  # ============================================================ #
  # YOUR CODE HERE:
  #   Calculate the output of the forward pass.  Notice the dimensions
  #   of w are D x M, which is the transpose of what we did in earlier
  #   assignments.
  # ============================================================ #

  # Copy and paste your code from HW #3
  N = x.shape[0]
  D = np.prod(x.shape[1:])
  x2 = np.reshape(x, (N, D))
  out = np.dot(x2, w) + b
  pass


  # ============================================================ #
  # END YOUR CODE HERE
  # ============================================================ #

  cache = (x, w, b)
  return out, cache
```

```
def affine_backward(dout, cache):
  """
  Computes the backward pass for an affine layer.

  Inputs:
  - dout: Upstream derivative, of shape (N, M)
  - cache: Tuple of:
    - x: Input data, of shape (N, d_1, ... d_k)
    - w: Weights, of shape (D, M)

  Returns a tuple of:
  - dx: Gradient with respect to x, of shape (N, d1, ..., d_k)
  - dw: Gradient with respect to w, of shape (D, M)
  - db: Gradient with respect to b, of shape (M,)
  """
  x, w, b = cache
  dx, dw, db = None, None, None

  # ============================================================= #
  # YOUR CODE HERE:
  #   Calculate the gradients for the backward pass.
  # ============================================================= #

  # Copy and paste your code from HW #3
  pass
  N = x.shape[0]
  D = np.prod(x.shape[1:])
  dx = np.dot(dout, w.T)
  dx = np.reshape(dx, x.shape)

  x2 = np.reshape(x, (N, D))
  dw = np.dot(x2.T, dout)
  db = np.sum(dout, axis = 0)

  # ============================================================= #
  # END YOUR CODE HERE
  # ============================================================= #

  return dx, dw, db

def relu_forward(x):
  """
  Computes the forward pass for a layer of rectified linear units (ReLUs).

  Input:
  - x: Inputs, of any shape

  Returns a tuple of:
  - out: Output, of the same shape as x
  - cache: x
  """
  # ============================================================= #
  # YOUR CODE HERE:
```

```
  #   Implement the ReLU forward pass.
  # ============================================================= #

  # Copy and paste your code from HW #3
  pass
  out = np.maximum(x, 0)

  # ============================================================= #
  # END YOUR CODE HERE
  # ============================================================= #

  cache = x
  return out, cache


def relu_backward(dout, cache):
  """
  Computes the backward pass for a layer of rectified linear units (ReLUs).

  Input:
  - dout: Upstream derivatives, of any shape
  - cache: Input x, of same shape as dout

  Returns:
  - dx: Gradient with respect to x
  """
  x = cache

  # ============================================================= #
  # YOUR CODE HERE:
  #   Implement the ReLU backward pass
  # ============================================================= #

  # Copy and paste your code from HW #3

  pass
  dx = np.array(dout, copy = True)
  dx[x<=0] = 0
#  out = np.maximum(0,x)
#  out[out>0] = 1
#  dx = out*dout

  # ============================================================= #
  # END YOUR CODE HERE
  # ============================================================= #

  return dx

def batchnorm_forward(x, gamma, beta, bn_param):
  """
  Forward pass for batch normalization.

  During training the sample mean and (uncorrected) sample variance are
  computed from minibatch statistics and used to normalize the incoming data.
```

```
During training we also keep an exponentially decaying running mean of the mean
and variance of each feature, and these averages are used to normalize data
at test-time.

At each timestep we update the running averages for mean and variance using
an exponential decay based on the momentum parameter:

running_mean = momentum * running_mean + (1 - momentum) * sample_mean
running_var = momentum * running_var + (1 - momentum) * sample_var

Note that the batch normalization paper suggests a different test-time
behavior: they compute sample mean and variance for each feature using a
large number of training images rather than using a running average. For
this implementation we have chosen to use running averages instead since
they do not require an additional estimation step; the torch7 implementation
of batch normalization also uses running averages.

Input:
- x: Data of shape (N, D)
- gamma: Scale parameter of shape (D,)
- beta: Shift paremeter of shape (D,)
- bn_param: Dictionary with the following keys:
  - mode: 'train' or 'test'; required
  - eps: Constant for numeric stability
  - momentum: Constant for running mean / variance.
  - running_mean: Array of shape (D,) giving running mean of features
  - running_var Array of shape (D,) giving running variance of features

Returns a tuple of:
- out: of shape (N, D)
- cache: A tuple of values needed in the backward pass
"""
mode = bn_param['mode']
eps = bn_param.get('eps', 1e-5)
momentum = bn_param.get('momentum', 0.9)

N, D = x.shape
running_mean = bn_param.get('running_mean', np.zeros(D, dtype=x.dtype))
running_var = bn_param.get('running_var', np.zeros(D, dtype=x.dtype))

out, cache = None, None
if mode == 'train':

  # ============================================================ #
  # YOUR CODE HERE:
  #   A few steps here:
  #     (1) Calculate the running mean and variance of the minibatch.
  #     (2) Normalize the activations with the running mean and variance.
  #     (3) Scale and shift the normalized activations.  Store this
  #         as the variable 'out'
  #     (4) Store any variables you may need for the backward pass in
  #         the 'cache' variable.
  # ============================================================ #
```

```
    pass
    mu = np.mean(x, axis = 0)
    var = np.var(x, axis = 0)
    x_hat = (x - mu)/np.sqrt(var + eps)
    out = gamma * x_hat + beta

    running_mean = momentum * running_mean + (1-momentum) * mu
    running_var = momentum * running_var + (1-momentum) * var
    cache = (x, out, beta, gamma, x_hat, mu, var, eps)

    # ============================================================ #
    # END YOUR CODE HERE
    # ============================================================ #

  elif mode == 'test':

    # ============================================================ #
    # YOUR CODE HERE:
    #   Calculate the testing time normalized activations.  Normalize using
    #   the running mean and variance, and then scale and shift appropriately.
    #   Store the output as 'out'.
    # ============================================================ #
    pass
    mu = running_mean
    var = running_var
    x_hat = (x - mu)/np.sqrt(var+eps)
    out = gamma*x_hat+beta
    cache = (mu, var, x_hat, out, gamma, beta)


    # ============================================================ #
    # END YOUR CODE HERE
    # ============================================================ #

  else:
    raise ValueError('Invalid forward batchnorm mode "%s"' % mode)

  # Store the updated running means back into bn_param
  bn_param['running_mean'] = running_mean
  bn_param['running_var'] = running_var

  return out, cache

def batchnorm_backward(dout, cache):
  """
  Backward pass for batch normalization.

  For this implementation, you should write out a computation graph for
  batch normalization on paper and propagate gradients backward through
  intermediate nodes.

  Inputs:
  - dout: Upstream derivatives, of shape (N, D)
  - cache: Variable of intermediates from batchnorm_forward.
```

```
  Returns a tuple of:
  - dx: Gradient with respect to inputs x, of shape (N, D)
  - dgamma: Gradient with respect to scale parameter gamma, of shape (D,)
  - dbeta: Gradient with respect to shift parameter beta, of shape (D,)
  """
  dx, dgamma, dbeta = None, None, None

  # ============================================================== #
  # YOUR CODE HERE:
  #   Implement the batchnorm backward pass, calculating dx, dgamma, and dbeta.
  # ============================================================== #
  x, out, beta, gamma, x_hat, mu, var, eps = cache
  N, D = x_hat.shape
  dbeta = np.sum(dout, axis = 0)  #(1, D)
  dgamma = np.sum(dout*x_hat, axis = 0) #(1, D)
  dx_hat = gamma*dout
  da = dx_hat/np.sqrt(var + eps)
  dmu = -np.sum(da, axis = 0)
  db = (x-mu) * dx_hat
  dc = -db/(var+eps)
  de = -0.5/(np.sqrt(var+eps)**3)*db
  dvar = np.sum(de, axis = 0)
  dx = dx_hat/np.sqrt(var+eps) + 2*(x-mu)*dvar/N + dmu/N

  # ============================================================== #
  # END YOUR CODE HERE
  # ============================================================== #

  return dx, dgamma, dbeta

def dropout_forward(x, dropout_param):
  """
  Performs the forward pass for (inverted) dropout.

  Inputs:
  - x: Input data, of any shape
  - dropout_param: A dictionary with the following keys:
    - p: Dropout parameter. We drop each neuron output with probability p.
    - mode: 'test' or 'train'. If the mode is train, then perform dropout;
      if the mode is test, then just return the input.
    - seed: Seed for the random number generator. Passing seed makes this
      function deterministic, which is needed for gradient checking but not in
      real networks.

  Outputs:
  - out: Array of the same shape as x.
  - cache: A tuple (dropout_param, mask). In training mode, mask is the dropout
    mask that was used to multiply the input; in test mode, mask is None.
  """
  p, mode = dropout_param['p'], dropout_param['mode']
  if 'seed' in dropout_param:
    np.random.seed(dropout_param['seed'])
```

```
  mask = None
  out = None

  if mode == 'train':
    # ============================================================ #
    # YOUR CODE HERE:
    #   Implement the inverted dropout forward pass during training time.
    #   Store the masked and scaled activations in out, and store the
    #   dropout mask as the variable mask.
    # ============================================================ #

    pass
    mask = (np.random.rand(*x.shape)<p)/p
    out = x*mask

    # ============================================================ #
    # END YOUR CODE HERE
    # ============================================================ #

  elif mode == 'test':

    # ============================================================ #
    # YOUR CODE HERE:
    #   Implement the inverted dropout forward pass during test time.
    # ============================================================ #

    pass
    mask = None
    out = x

    # ============================================================ #
    # END YOUR CODE HERE
    # ============================================================ #

  cache = (dropout_param, mask)
  out = out.astype(x.dtype, copy=False)

  return out, cache

def dropout_backward(dout, cache):
  """
  Perform the backward pass for (inverted) dropout.

  Inputs:
  - dout: Upstream derivatives, of any shape
  - cache: (dropout_param, mask) from dropout_forward.
  """
  dropout_param, mask = cache
  mode = dropout_param['mode']

  dx = None
  if mode == 'train':
    # ============================================================ #
    # YOUR CODE HERE:
```

```
    #   Implement the inverted dropout backward pass during training time.
    # ============================================================ #
    pass
    dx = dout*mask
    # ============================================================ #
    # END YOUR CODE HERE
    # ============================================================ #
  elif mode == 'test':
    # ============================================================ #
    # YOUR CODE HERE:
    #   Implement the inverted dropout backward pass during test time.
    # ============================================================ #
    pass
    dx = dout
    # ============================================================ #
    # END YOUR CODE HERE
    # ============================================================ #
  return dx

def svm_loss(x, y):
  """
  Computes the loss and gradient using for multiclass SVM classification.

  Inputs:
  - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
    for the ith input.
  - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
    0 <= y[i] < C

  Returns a tuple of:
  - loss: Scalar giving the loss
  - dx: Gradient of the loss with respect to x
  """
  N = x.shape[0]
  correct_class_scores = x[np.arange(N), y]
  margins = np.maximum(0, x - correct_class_scores[:, np.newaxis] + 1.0)
  margins[np.arange(N), y] = 0
  loss = np.sum(margins) / N
  num_pos = np.sum(margins > 0, axis=1)
  dx = np.zeros_like(x)
  dx[margins > 0] = 1
  dx[np.arange(N), y] -= num_pos
  dx /= N
  return loss, dx


def softmax_loss(x, y):
  """
  Computes the loss and gradient for softmax classification.

  Inputs:
  - x: Input data, of shape (N, C) where x[i, j] is the score for the jth class
    for the ith input.
  - y: Vector of labels, of shape (N,) where y[i] is the label for x[i] and
```

```
  0 <= y[i] < C

Returns a tuple of:
- loss: Scalar giving the loss
- dx: Gradient of the loss with respect to x
"""

probs = np.exp(x - np.max(x, axis=1, keepdims=True))
probs /= np.sum(probs, axis=1, keepdims=True)
N = x.shape[0]
loss = -np.sum(np.log(probs[np.arange(N), y]+1e-7)) / N
dx = probs.copy()
dx[np.arange(N), y] -= 1
dx /= N
return loss, dx
```

```
from .layers import *

"""
This code was originally written for CS 231n at Stanford University
(cs231n.stanford.edu).  It has been modified in various areas for use in the
ECE 239AS class at UCLA.  This includes the descriptions of what code to
implement as well as some slight potential changes in variable names to be
consistent with class nomenclature.  We thank Justin Johnson & Serena Yeung for
permission to use this code.  To see the original version, please visit
cs231n.stanford.edu.
"""

def affine_relu_forward(x, w, b):
  """
  Convenience layer that performs an affine transform followed by a ReLU

  Inputs:
  - x: Input to the affine layer
  - w, b: Weights for the affine layer

  Returns a tuple of:
  - out: Output from the ReLU
  - cache: Object to give to the backward pass
  """
  a, fc_cache = affine_forward(x, w, b)
  out, relu_cache = relu_forward(a)
  cache = (fc_cache, relu_cache)
  return out, cache


def affine_relu_backward(dout, cache):
  """
  Backward pass for the affine-relu convenience layer
  """
  fc_cache, relu_cache = cache
  da = relu_backward(dout, relu_cache)
  dx, dw, db = affine_backward(da, fc_cache)
  return dx, dw, db

def affine_batchnorm_relu_forward(x, w, b, gamma, beta, bn_param): #in a pass
 computation in loss
  aff_out, fc_cache = affine_forward(x, w, b)
  bn_out, bn_cache = batchnorm_forward(aff_out, gamma, beta, bn_param)
  out, relu_cache = relu_forward(bn_out)
  cache = (fc_cache, bn_cache, relu_cache)
  return out, cache

def affine_batchnorm_relu_backward(dout, cache):
  fc_cache, bn_cache, relu_cache = cache
  da = relu_backward(dout, relu_cache)
  dx_norm, dgamma_norm, dbeta_norm = batchnorm_backward(da, bn_cache)
  dx, dw, db = affine_backward(dx_norm, fc_cache)
  return dx, dw, db, dgamma_norm, dbeta_norm
```

```
import numpy as np
import pdb

from .layers import *
from .layer_utils import *

"""
This code was originally written for CS 231n at Stanford University
(cs231n.stanford.edu).  It has been modified in various areas for use in the
ECE 239AS class at UCLA.  This includes the descriptions of what code to
implement as well as some slight potential changes in variable names to be
consistent with class nomenclature.  We thank Justin Johnson & Serena Yeung for
permission to use this code.  To see the original version, please visit
cs231n.stanford.edu.
"""

class FullyConnectedNet(object):
  """
  A fully-connected neural network with an arbitrary number of hidden layers,
  ReLU nonlinearities, and a softmax loss function. This will also implement
  dropout and batch normalization as options. For a network with L layers,
  the architecture will be

  {affine - [batch norm] - relu - [dropout]} x (L - 1) - affine - softmax

  where batch normalization and dropout are optional, and the {...} block is
  repeated L - 1 times.

  Similar to the TwoLayerNet above, learnable parameters are stored in the
  self.params dictionary and will be learned using the Solver class.
  """

  def __init__(self, hidden_dims, input_dim=3*32*32, num_classes=10,
               dropout=0, use_batchnorm=False, reg=0.0,
               weight_scale=1e-2, dtype=np.float32, seed=None):
    """
    Initialize a new FullyConnectedNet.

    Inputs:
    - hidden_dims: A list of integers giving the size of each hidden layer.
    - input_dim: An integer giving the size of the input.
    - num_classes: An integer giving the number of classes to classify.
    - dropout: Scalar between 0 and 1 giving dropout strength. If dropout=0 then
      the network should not use dropout at all.
    - use_batchnorm: Whether or not the network should use batch normalization.
    - reg: Scalar giving L2 regularization strength.
    - weight_scale: Scalar giving the standard deviation for random
      initialization of the weights.
    - dtype: A numpy datatype object; all computations will be performed using
      this datatype. float32 is faster but less accurate, so you should use
      float64 for numeric gradient checking.
    - seed: If not None, then pass this random seed to the dropout layers. This
      will make the dropout layers deteriminstic so we can gradient check the
      model.
```

```
"""
self.use_batchnorm = use_batchnorm
self.use_dropout = dropout > 0
self.reg = reg
self.num_layers = 1 + len(hidden_dims)
self.dtype = dtype
self.params = {}

# ============================================================ #
# YOUR CODE HERE:
#   Initialize all parameters of the network in the self.params dictionary.
#   The weights and biases of layer 1 are W1 and b1; and in general the
#   weights and biases of layer i are Wi and bi. The
#   biases are initialized to zero and the weights are initialized
#   so that each parameter has mean 0 and standard deviation weight_scale.
#
#   BATCHNORM: Initialize the gammas of each layer to 1 and the beta
#   parameters to zero.  The gamma and beta parameters for layer 1 should
#   be self.params['gamma1'] and self.params['beta1'].  For layer 2, they
#   should be gamma2 and beta2, etc. Only use batchnorm if self.use_batchnorm
#   is true and DO NOT batch normalize the output scores.
# ============================================================ #
self.h_dims = len(hidden_dims)
for i in np.arange(len(hidden_dims) + 1):
    # eg: when len(hidden_dims) = 2, W1, W2, W3, b1, b2, b3 exist
    if i == 0:
        self.params['W'+str(i+1)] = weight_scale * np.random.randn(input_dim,
         hidden_dims[i])
        self.params['b'+str(i+1)] = np.zeros(hidden_dims[i])
    elif i == len(hidden_dims):
        self.params['W'+str(i+1)] = weight_scale *
         np.random.randn(hidden_dims[-1], num_classes)
        self.params['b'+str(i+1)] = np.zeros(num_classes)
    else:
        self.params['W'+str(i+1)] = weight_scale *
         np.random.randn(hidden_dims[i-1], hidden_dims[i])
        self.params['b'+str(i+1)] = np.zeros(hidden_dims[i])

pass

# ============================================================ #
# END YOUR CODE HERE
# ============================================================ #

# When using dropout we need to pass a dropout_param dictionary to each
# dropout layer so that the layer knows the dropout probability and the mode
# (train / test). You can pass the same dropout_param to each dropout layer.
self.dropout_param = {}
if self.use_dropout:
  self.dropout_param = {'mode': 'train', 'p': dropout}
  if seed is not None:
    self.dropout_param['seed'] = seed

# With batch normalization we need to keep track of running means and
```

```
    # variances, so we need to pass a special bn_param object to each batch
    # normalization layer. You should pass self.bn_params[0] to the forward pass
    # of the first batch normalization layer, self.bn_params[1] to the forward
    # pass of the second batch normalization layer, etc.
    self.bn_params = []
    if self.use_batchnorm:
        self.bn_params = [{'mode': 'train'} for i in np.arange(self.num_layers - 1)]
        # initialize all gammas and betas
        for i in np.arange(self.num_layers-1):
            self.params['gamma'+str(i+1)] = np.ones(hidden_dims[i])
            self.params['beta'+str(i+1)] = np.zeros(hidden_dims[i])


    # Cast all parameters to the correct datatype
    for k, v in self.params.items():
      self.params[k] = v.astype(dtype)


  def loss(self, X, y=None):
    """
    Compute loss and gradient for the fully-connected net.

    Input / output: Same as TwoLayerNet above.
    """
    X = X.astype(self.dtype)
    mode = 'test' if y is None else 'train'

    # Set train/test mode for batchnorm params and dropout param since they
    # behave differently during training and testing.
    if self.dropout_param is not None:
      self.dropout_param['mode'] = mode
    if self.use_batchnorm:
      for bn_param in self.bn_params:
        bn_param[mode] = mode

    scores = None

    # ============================================================== #
    # YOUR CODE HERE:
    #   Implement the forward pass of the FC net and store the output
    #   scores as the variable "scores".
    #
    #   BATCHNORM: If self.use_batchnorm is true, insert a bathnorm layer
    #   between the affine_forward and relu_forward layers.  You may
    #   also write an affine_batchnorm_relu() function in layer_utils.py.
    #
    #   DROPOUT: If dropout is non-zero, insert a dropout layer after
    #   every ReLU layer.
    # ============================================================== #
    N = X.shape[0]
    D = np.prod(X.shape[1:])

    hidden = {}
    hidden['h0'] = X
```

```
    h = X

    if self.use_dropout:
        hidden['hdrop0'], hidden['hdrop_cache0'] = dropout_forward(h,
         self.dropout_param)
        h = hidden['hdrop0']

    for i in np.arange(self.h_dims + 1):
        w = self.params['W'+str(i+1)]
        b = self.params['b'+str(i+1)]

        if i == self.h_dims:
            scores, scores_cache = affine_forward(h, w, b)
            hidden['h'+str(i+1)] = h
            hidden['h_cache'+str(i+1)] = h_cache
        else:
            if self.use_batchnorm == True:
                gamma = self.params['gamma'+str(i+1)]
                beta = self.params['beta'+str(i+1)]
                bn_param = self.bn_params[i]
                h, h_cache = affine_batchnorm_relu_forward(h, w, b, gamma, beta,
                 bn_param)
                hidden['h'+str(i+1)] = h
                hidden['h_cache'+str(i+1)] = h_cache
            else:
                h, h_cache = affine_relu_forward(h, w, b)
                hidden['h'+str(i+1)] = h
                hidden['h_cache'+str(i+1)] = h_cache

            if self.use_dropout:
                hidden['hdrop'+str(i+1)], hidden['hdrop_cache'+str(i+1)] =
                 dropout_forward(h, self.dropout_param)
                h = hidden['hdrop'+str(i+1)]


#         if i == 0:
#             h, h_cache = affine_relu_forward(X, w, b)
#             hidden['h'+str(i+1)] = h
#             hidden['h_cache'+str(i+1)] = h_cache
#         elif i == self.h_dims:
#             scores, scores_cache = affine_forward(h, w, b)
#             hidden['h'+str(i+1)] = h
#             hidden['h_cache'+str(i+1)] = h_cache
#         else:
#             h, h_cache = affine_relu_forward(h, w, b)
#             hidden['h'+str(i+1)] = h
#             hidden['h_cache'+str(i+1)] = h_cache
    pass

    # ============================================================ #
    # END YOUR CODE HERE
    # ============================================================ #

    # If test mode return early
```

```
if mode == 'test':
  return scores

loss, grads = 0.0, {}

# ============================================================ #
# YOUR CODE HERE:
#   Implement the backwards pass of the FC net and store the gradients
#   in the grads dict, so that grads[k] is the gradient of self.params[k]
#   Be sure your L2 regularization includes a 0.5 factor.
#
#   BATCHNORM: Incorporate the backward pass of the batchnorm.
#
#   DROPOUT: Incorporate the backward pass of dropout.
# ============================================================ #
data_loss, dscores = softmax_loss(scores, y)
reg_sum = 0
for i in np.arange(self.h_dims+1):
    reg_sum += np.sum(self.params['W'+str(i+1)]*self.params['W'+str(i+1)])
reg_loss = 0.5 * self.reg * reg_sum
loss = data_loss + reg_loss

for i in np.arange(self.h_dims+1)[::-1]:
    if i == self.h_dims:
        dloss, w, b = affine_backward(dscores, scores_cache)
        grads['W'+str(i+1)] = w
        grads['W'+str(i+1)] += self.reg * self.params['W'+str(i+1)]
        grads['b'+str(i+1)] = b
    else:
        if self.use_dropout:
            dloss  = dropout_backward(dloss, hidden['hdrop_cache'+str(i+1)])

        if self.use_batchnorm == True:
            dloss, dw, db, dgamma, dbeta = affine_batchnorm_relu_backward(dloss,
             hidden['h_cache'+str(i+1)])
            grads['W'+str(i+1)] = dw
            grads['b'+str(i+1)] = db
            grads['W'+str(i+1)] += self.reg * self.params['W'+str(i+1)]
            grads['gamma'+str(i+1)] = dgamma
            grads['beta'+str(i+1)] = dbeta
        else:
            dloss, dw, db = affine_relu_backward(dloss,
             hidden['h_cache'+str(i+1)])
            grads['W'+str(i+1)] = dw
            grads['b'+str(i+1)] = db
            grads['W'+str(i+1)] += self.reg * self.params['W'+str(i+1)]


pass
# ============================================================ #
# END YOUR CODE HERE
# ============================================================ #

return loss, grads
```

Dropout

In this notebook, you will implement dropout. Then we will ask you to train a network with batchnorm and dropout, and acheive over 60% accuracy on CIFAR-10.

CS231n has built a solid API for building these modular frameworks and training them, and we will use their very well implemented framework as opposed to "reinventing the wheel." This includes using their Solver, various utility functions, and their layer structure. This also includes nndl.fc_net, nndl.layers, and nndl.layer_utils. As in prior assignments, we thank Serena Yeung & Justin Johnson for permission to use code written for the CS 231n class (cs231n.stanford.edu).

In [1]:

```
## Import and setups

import time
import numpy as np
import matplotlib.pyplot as plt
from nndl.fc_net import *
from nndl.layers import *
from cs231n.data_utils import get_CIFAR10_data
from cs231n.gradient_check import eval_numerical_gradient, eval_numerical_gradient_a
from cs231n.solver import Solver

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading external modules
# see http://stackoverflow.com/questions/1907993/autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2

def rel_error(x, y):
  """ returns relative error """
  return np.max(np.abs(x - y) / (np.maximum(1e-8, np.abs(x) + np.abs(y))))
```

In [2]:

```
# Load the (preprocessed) CIFAR10 data.

data = get_CIFAR10_data()
for k in data.keys():
  print('{}: {} '.format(k, data[k].shape))
```

```
X_train: (49000, 3, 32, 32)
y_train: (49000,)
X_val: (1000, 3, 32, 32)
y_val: (1000,)
X_test: (1000, 3, 32, 32)
y_test: (1000,)
```

Dropout forward pass

Implement the training and test time dropout forward pass, `dropout_forward` , in `nndl/layers.py` . After that, test your implementation by running the following cell.

In [3]:

```
x = np.random.randn(500, 500) + 10

for p in [0.3, 0.6, 0.75]:
  out, _ = dropout_forward(x, {'mode': 'train', 'p': p})
  out_test, _ = dropout_forward(x, {'mode': 'test', 'p': p})

  print('Running tests with p = ', p)
  print('Mean of input: ', x.mean())
  print('Mean of train-time output: ', out.mean())
  print('Mean of test-time output: ', out_test.mean())
  print('Fraction of train-time output set to zero: ', (out == 0).mean())
  print('Fraction of test-time output set to zero: ', (out_test == 0).mean())
```

```
Running tests with p =  0.3
Mean of input:  9.999173896404256
Mean of train-time output:  10.03564461966431
Mean of test-time output:  9.999173896404256
Fraction of train-time output set to zero:  0.69882
Fraction of test-time output set to zero:  0.0
Running tests with p =  0.6
Mean of input:  9.999173896404256
Mean of train-time output:  9.981303066021745
Mean of test-time output:  9.999173896404256
Fraction of train-time output set to zero:  0.401032
Fraction of test-time output set to zero:  0.0
Running tests with p =  0.75
Mean of input:  9.999173896404256
Mean of train-time output:  10.001870004386827
Mean of test-time output:  9.999173896404256
Fraction of train-time output set to zero:  0.249708
Fraction of test-time output set to zero:  0.0
```

Dropout backward pass

Implement the backward pass, `dropout_backward` , in `nndl/layers.py` . After that, test your gradients by running the following cell:

In [4]:

```
x = np.random.randn(10, 10) + 10
dout = np.random.randn(*x.shape)

dropout_param = {'mode': 'train', 'p': 0.8, 'seed': 123}
out, cache = dropout_forward(x, dropout_param)
dx = dropout_backward(dout, cache)
dx_num = eval_numerical_gradient_array(lambda xx: dropout_forward(xx, dropout_param)

print('dx relative error: ', rel_error(dx, dx_num))
```

```
dx relative error:  5.445613574740722e-11
```

Implement a fully connected neural network with dropout layers

Modify the `FullyConnectedNet()` class in `nndl/fc_net.py` to incorporate dropout. A dropout layer should be incorporated after every ReLU layer. Concretely, there shouldn't be a dropout at the output layer since there is no ReLU at the output layer. You will need to modify the class in the following areas:

(1) In the forward pass, you will need to incorporate a dropout layer after every relu layer.

(2) In the backward pass, you will need to incorporate a dropout backward pass layer.

Check your implementation by running the following code. Our W1 gradient relative error is on the order of 1e-6 (the largest of all the relative errors).

Dropout as a regularizer

In class, we claimed that dropout acts as a regularizer by effectively bagging. To check this, we will train two small networks, one with dropout and one without dropout.

In [5]:

```
N, D, H1, H2, C = 2, 15, 20, 30, 10
X = np.random.randn(N, D)
y = np.random.randint(C, size=(N,))

for dropout in [0, 0.25, 0.5]:
  print('Running check with dropout = ', dropout)
  model = FullyConnectedNet([H1, H2], input_dim=D, num_classes=C,
                            weight_scale=5e-2, dtype=np.float64,
                            dropout=dropout, seed=123)

  loss, grads = model.loss(X, y)
  print('Initial loss: ', loss)

  for name in sorted(grads):
    f = lambda _: model.loss(X, y)[0]
    grad_num = eval_numerical_gradient(f, model.params[name], verbose=False, h=1e-5)
    print('{} relative error: {}'.format(name, rel_error(grad_num, grads[name])))
  print('\n')
```

```
Running check with dropout =  0
Initial loss:  2.3051938247860218
W1 relative error: 7.711419522175973e-07
W2 relative error: 1.5034484932141387e-05
W3 relative error: 6.398873279808276e-07
b1 relative error: 2.9369574464090924e-06
b2 relative error: 6.320762652162454e-07
b3 relative error: 5.016029374797846e-07


Running check with dropout =  0.25
Initial loss:  2.305429202994099
W1 relative error: 5.024649969461409e-07
W2 relative error: 5.026393983338925e-07
W3 relative error: 5.292126245232686e-07
b1 relative error: 5.029711099712491e-07
b2 relative error: 5.02533626843244e-07
b3 relative error: 5.026138778638627e-07


Running check with dropout =  0.5
Initial loss:  2.307550515805681
W1 relative error: 5.399786757971378e-07
W2 relative error: 5.345642293235742e-07
W3 relative error: 5.440162773705216e-07
b1 relative error: 5.139377993145615e-07
b2 relative error: 5.056039855824707e-07
b3 relative error: 5.060262454954197e-07
```

In [6]:

```
# Train two identical nets, one with dropout and one without

num_train = 500
small_data = {
  'X_train': data['X_train'][:num_train],
  'y_train': data['y_train'][:num_train],
  'X_val': data['X_val'],
  'y_val': data['y_val'],
}

solvers = {}
dropout_choices = [0, 0.6]
for dropout in dropout_choices:
  model = FullyConnectedNet([100, 100, 100], dropout=dropout)

  solver = Solver(model, small_data,
                  num_epochs=25, batch_size=100,
                  update_rule='adam',
                  optim_config={
                    'learning_rate': 5e-4,
                  },
                  verbose=True, print_every=100)
  solver.train()
  solvers[dropout] = solver
```

```
(Iteration 1 / 125) loss: 2.300804
(Epoch 0 / 25) train acc: 0.220000; val_acc: 0.168000
(Epoch 1 / 25) train acc: 0.188000; val_acc: 0.147000
(Epoch 2 / 25) train acc: 0.266000; val_acc: 0.200000
(Epoch 3 / 25) train acc: 0.338000; val_acc: 0.262000
(Epoch 4 / 25) train acc: 0.378000; val_acc: 0.278000
(Epoch 5 / 25) train acc: 0.428000; val_acc: 0.297000
(Epoch 6 / 25) train acc: 0.468000; val_acc: 0.323000
(Epoch 7 / 25) train acc: 0.494000; val_acc: 0.287000
(Epoch 8 / 25) train acc: 0.566000; val_acc: 0.328000
(Epoch 9 / 25) train acc: 0.572000; val_acc: 0.322000
(Epoch 10 / 25) train acc: 0.622000; val_acc: 0.324000
(Epoch 11 / 25) train acc: 0.670000; val_acc: 0.279000
(Epoch 12 / 25) train acc: 0.710000; val_acc: 0.338000
(Epoch 13 / 25) train acc: 0.746000; val_acc: 0.319000
(Epoch 14 / 25) train acc: 0.792000; val_acc: 0.307000
(Epoch 15 / 25) train acc: 0.834000; val_acc: 0.297000
(Epoch 16 / 25) train acc: 0.876000; val_acc: 0.327000
(Epoch 17 / 25) train acc: 0.886000; val_acc: 0.320000
(Epoch 18 / 25) train acc: 0.918000; val_acc: 0.314000
(Epoch 19 / 25) train acc: 0.922000; val_acc: 0.290000
(Epoch 20 / 25) train acc: 0.944000; val_acc: 0.306000
(Iteration 101 / 125) loss: 0.156105
(Epoch 21 / 25) train acc: 0.968000; val_acc: 0.302000
(Epoch 22 / 25) train acc: 0.978000; val_acc: 0.302000
(Epoch 23 / 25) train acc: 0.976000; val_acc: 0.289000
(Epoch 24 / 25) train acc: 0.986000; val_acc: 0.285000
(Epoch 25 / 25) train acc: 0.978000; val_acc: 0.311000
```

```
(Iteration 1 / 125) loss: 2.309402
(Epoch 0 / 25) train acc: 0.130000; val_acc: 0.151000
(Epoch 1 / 25) train acc: 0.220000; val_acc: 0.168000
(Epoch 2 / 25) train acc: 0.232000; val_acc: 0.188000
(Epoch 3 / 25) train acc: 0.220000; val_acc: 0.180000
(Epoch 4 / 25) train acc: 0.236000; val_acc: 0.214000
(Epoch 5 / 25) train acc: 0.306000; val_acc: 0.258000
(Epoch 6 / 25) train acc: 0.308000; val_acc: 0.269000
(Epoch 7 / 25) train acc: 0.344000; val_acc: 0.290000
(Epoch 8 / 25) train acc: 0.358000; val_acc: 0.285000
(Epoch 9 / 25) train acc: 0.396000; val_acc: 0.305000
(Epoch 10 / 25) train acc: 0.406000; val_acc: 0.307000
(Epoch 11 / 25) train acc: 0.428000; val_acc: 0.319000
(Epoch 12 / 25) train acc: 0.458000; val_acc: 0.321000
(Epoch 13 / 25) train acc: 0.486000; val_acc: 0.328000
(Epoch 14 / 25) train acc: 0.490000; val_acc: 0.326000
(Epoch 15 / 25) train acc: 0.540000; val_acc: 0.334000
(Epoch 16 / 25) train acc: 0.540000; val_acc: 0.328000
(Epoch 17 / 25) train acc: 0.562000; val_acc: 0.321000
(Epoch 18 / 25) train acc: 0.590000; val_acc: 0.326000
(Epoch 19 / 25) train acc: 0.604000; val_acc: 0.324000
(Epoch 20 / 25) train acc: 0.600000; val_acc: 0.322000
(Iteration 101 / 125) loss: 1.377055
(Epoch 21 / 25) train acc: 0.628000; val_acc: 0.321000
(Epoch 22 / 25) train acc: 0.644000; val_acc: 0.295000
(Epoch 23 / 25) train acc: 0.622000; val_acc: 0.330000
(Epoch 24 / 25) train acc: 0.650000; val_acc: 0.340000
(Epoch 25 / 25) train acc: 0.708000; val_acc: 0.330000
```

In [7]:

```
# Plot train and validation accuracies of the two models

train_accs = []
val_accs = []
for dropout in dropout_choices:
  solver = solvers[dropout]
  train_accs.append(solver.train_acc_history[-1])
  val_accs.append(solver.val_acc_history[-1])

plt.subplot(3, 1, 1)
for dropout in dropout_choices:
  plt.plot(solvers[dropout].train_acc_history, 'o', label='%.2f dropout' % dropout)
plt.title('Train accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(ncol=2, loc='lower right')

plt.subplot(3, 1, 2)
for dropout in dropout_choices:
  plt.plot(solvers[dropout].val_acc_history, 'o', label='%.2f dropout' % dropout)
plt.title('Val accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend(ncol=2, loc='lower right')

plt.gcf().set_size_inches(15, 15)
plt.show()
```

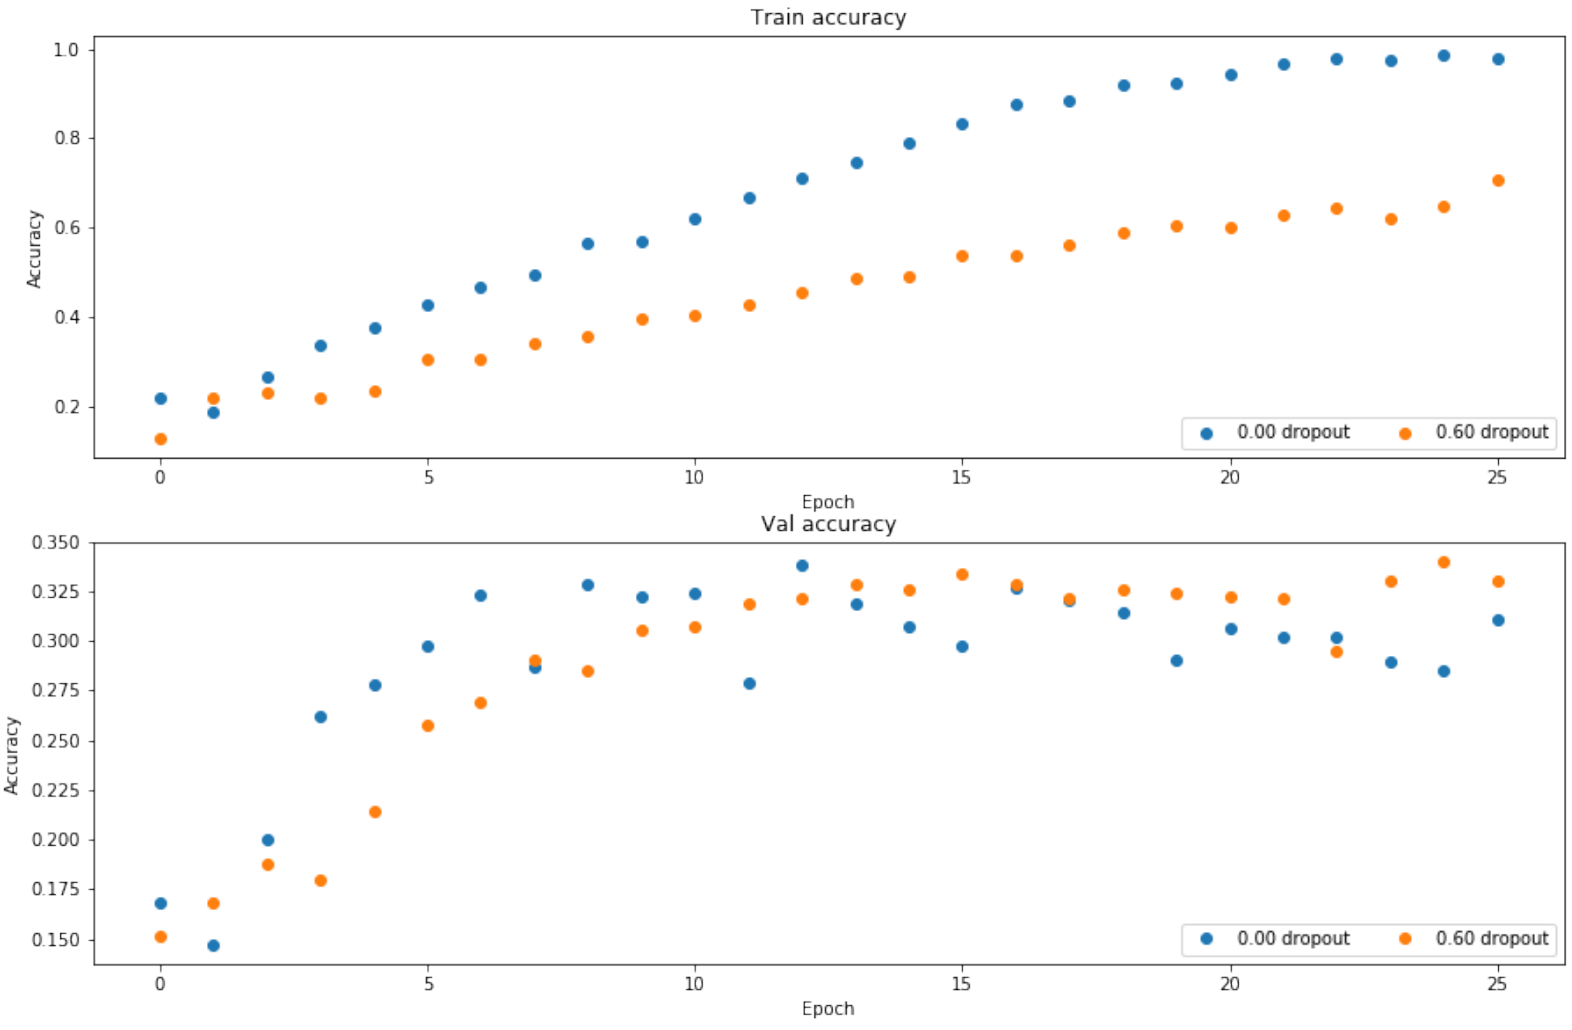

Question

Based off the results of this experiment, is dropout performing regularization? Explain your answer.

Answer:

Yes. It can be seen that by using dropout, the training accuracy decreases while the validation accuracy increases. Dropout can help reduce the overfitting problem so that is performing regularization.

Final part of the assignment

Get over 60% validation accuracy on CIFAR-10 by using the layers you have implemented. You will be graded according to the following equation:

min(floor((X - 32%)) / 28%, 1) where if you get 60% or higher validation accuracy, you get full points.

In [8]:

```
# ================================================================ #
# YOUR CODE HERE:
#   Implement a FC-net that achieves at least 60% validation accuracy
#   on CIFAR-10.
# ================================================================ #
data = get_CIFAR10_data()
hidden_dims = [500, 500, 500, 500]
whole_data = {
  'X_train': data['X_train'],
  'y_train': data['y_train'],
  'X_val': data['X_val'],
  'y_val': data['y_val'],
  'X_test': data['X_test'],
  'y_test': data['y_test'],

}
weight_scale = 1e-3 #
dropout = 0.7
model = FullyConnectedNet(hidden_dims, weight_scale=weight_scale, reg = 0, dropout =
solver = Solver(model, whole_data,
                num_epochs=60, batch_size=100,
                update_rule='adam',
                optim_config={'learning_rate': 1e-3,},
                lr_decay = 0.95,
                verbose=True, print_every=200)
solver.train()
# ================================================================ #
# END YOUR CODE HERE
# ================================================================ #
```

```
(Iteration 1 / 29400) loss: 2.298408
(Epoch 0 / 60) train acc: 0.141000; val_acc: 0.151000
(Iteration 201 / 29400) loss: 1.594431
(Iteration 401 / 29400) loss: 1.598612
(Epoch 1 / 60) train acc: 0.463000; val_acc: 0.448000
(Iteration 601 / 29400) loss: 1.723591
(Iteration 801 / 29400) loss: 1.557372
(Epoch 2 / 60) train acc: 0.462000; val_acc: 0.488000
(Iteration 1001 / 29400) loss: 1.598769
(Iteration 1201 / 29400) loss: 1.491483
(Iteration 1401 / 29400) loss: 1.463362
(Epoch 3 / 60) train acc: 0.521000; val_acc: 0.494000
(Iteration 1601 / 29400) loss: 1.509252
(Iteration 1801 / 29400) loss: 1.466622
(Epoch 4 / 60) train acc: 0.545000; val_acc: 0.529000
(Iteration 2001 / 29400) loss: 1.259682
(Iteration 2201 / 29400) loss: 1.257530
(Iteration 2401 / 29400) loss: 1.242700
(Epoch 5 / 60) train acc: 0.578000; val_acc: 0.550000
(Iteration 2601 / 29400) loss: 1.494734
(Iteration 2801 / 29400) loss: 1.319884
(Epoch 6 / 60) train acc: 0.574000; val_acc: 0.548000
(Iteration 3001 / 29400) loss: 1.412498
(Iteration 3201 / 29400) loss: 1.288664
(Iteration 3401 / 29400) loss: 1.226284
(Epoch 7 / 60) train acc: 0.611000; val_acc: 0.556000
(Iteration 3601 / 29400) loss: 1.143371
(Iteration 3801 / 29400) loss: 1.251719
(Epoch 8 / 60) train acc: 0.623000; val_acc: 0.555000
(Iteration 4001 / 29400) loss: 1.249114
(Iteration 4201 / 29400) loss: 1.172773
(Iteration 4401 / 29400) loss: 1.119578
(Epoch 9 / 60) train acc: 0.633000; val_acc: 0.559000
(Iteration 4601 / 29400) loss: 1.229059
(Iteration 4801 / 29400) loss: 1.262048
(Epoch 10 / 60) train acc: 0.636000; val_acc: 0.564000
(Iteration 5001 / 29400) loss: 1.354756
(Iteration 5201 / 29400) loss: 1.420874
(Epoch 11 / 60) train acc: 0.641000; val_acc: 0.569000
(Iteration 5401 / 29400) loss: 1.230693
(Iteration 5601 / 29400) loss: 1.043646
(Iteration 5801 / 29400) loss: 1.024204
(Epoch 12 / 60) train acc: 0.641000; val_acc: 0.581000
(Iteration 6001 / 29400) loss: 1.121003
(Iteration 6201 / 29400) loss: 1.223647
(Epoch 13 / 60) train acc: 0.633000; val_acc: 0.560000
(Iteration 6401 / 29400) loss: 0.978449
(Iteration 6601 / 29400) loss: 1.315928
(Iteration 6801 / 29400) loss: 1.060285
(Epoch 14 / 60) train acc: 0.656000; val_acc: 0.575000
(Iteration 7001 / 29400) loss: 1.254175
(Iteration 7201 / 29400) loss: 1.071293
(Epoch 15 / 60) train acc: 0.703000; val_acc: 0.575000
```

```
(Iteration 7401 / 29400) loss: 1.051179
(Iteration 7601 / 29400) loss: 0.993687
(Iteration 7801 / 29400) loss: 1.239175
(Epoch 16 / 60) train acc: 0.661000; val_acc: 0.580000
(Iteration 8001 / 29400) loss: 1.368725
(Iteration 8201 / 29400) loss: 0.893911
(Epoch 17 / 60) train acc: 0.725000; val_acc: 0.568000
(Iteration 8401 / 29400) loss: 0.965442
(Iteration 8601 / 29400) loss: 1.009901
(Iteration 8801 / 29400) loss: 1.008127
(Epoch 18 / 60) train acc: 0.705000; val_acc: 0.583000
(Iteration 9001 / 29400) loss: 0.983968
(Iteration 9201 / 29400) loss: 1.203774
(Epoch 19 / 60) train acc: 0.744000; val_acc: 0.582000
(Iteration 9401 / 29400) loss: 1.004626
(Iteration 9601 / 29400) loss: 1.013244
(Epoch 20 / 60) train acc: 0.716000; val_acc: 0.588000
(Iteration 9801 / 29400) loss: 1.084729
(Iteration 10001 / 29400) loss: 0.885756
(Iteration 10201 / 29400) loss: 1.114381
(Epoch 21 / 60) train acc: 0.726000; val_acc: 0.592000
(Iteration 10401 / 29400) loss: 1.118271
(Iteration 10601 / 29400) loss: 1.087094
(Epoch 22 / 60) train acc: 0.730000; val_acc: 0.600000
(Iteration 10801 / 29400) loss: 1.047215
(Iteration 11001 / 29400) loss: 1.079156
(Iteration 11201 / 29400) loss: 1.115825
(Epoch 23 / 60) train acc: 0.761000; val_acc: 0.591000
(Iteration 11401 / 29400) loss: 0.987570
(Iteration 11601 / 29400) loss: 1.156316
(Epoch 24 / 60) train acc: 0.748000; val_acc: 0.587000
(Iteration 11801 / 29400) loss: 0.899250

(Iteration 12001 / 29400) loss: 0.930228
(Iteration 12201 / 29400) loss: 0.869270
(Epoch 25 / 60) train acc: 0.748000; val_acc: 0.607000
(Iteration 12401 / 29400) loss: 0.982034
(Iteration 12601 / 29400) loss: 0.998646
(Epoch 26 / 60) train acc: 0.783000; val_acc: 0.591000
(Iteration 12801 / 29400) loss: 0.799769
(Iteration 13001 / 29400) loss: 1.040560
(Iteration 13201 / 29400) loss: 0.965084
(Epoch 27 / 60) train acc: 0.779000; val_acc: 0.589000
(Iteration 13401 / 29400) loss: 0.842854
(Iteration 13601 / 29400) loss: 1.041773
(Epoch 28 / 60) train acc: 0.741000; val_acc: 0.580000
(Iteration 13801 / 29400) loss: 0.942940
(Iteration 14001 / 29400) loss: 1.035560
(Iteration 14201 / 29400) loss: 0.768751
(Epoch 29 / 60) train acc: 0.790000; val_acc: 0.589000
(Iteration 14401 / 29400) loss: 0.847572
(Iteration 14601 / 29400) loss: 0.873548
(Epoch 30 / 60) train acc: 0.798000; val_acc: 0.582000
(Iteration 14801 / 29400) loss: 0.948969
```

```
(Iteration 15001 / 29400) loss: 0.902381
(Epoch 31 / 60) train acc: 0.793000; val_acc: 0.592000
(Iteration 15201 / 29400) loss: 0.772018
(Iteration 15401 / 29400) loss: 0.882011
(Iteration 15601 / 29400) loss: 0.807604
(Epoch 32 / 60) train acc: 0.786000; val_acc: 0.591000
(Iteration 15801 / 29400) loss: 0.765125
(Iteration 16001 / 29400) loss: 0.786065
(Epoch 33 / 60) train acc: 0.823000; val_acc: 0.582000
(Iteration 16201 / 29400) loss: 0.905458
(Iteration 16401 / 29400) loss: 0.811378
(Iteration 16601 / 29400) loss: 0.824763
(Epoch 34 / 60) train acc: 0.792000; val_acc: 0.601000
(Iteration 16801 / 29400) loss: 0.920568
(Iteration 17001 / 29400) loss: 0.926238
(Epoch 35 / 60) train acc: 0.794000; val_acc: 0.592000
(Iteration 17201 / 29400) loss: 0.840497
(Iteration 17401 / 29400) loss: 0.818339
(Iteration 17601 / 29400) loss: 1.023683
(Epoch 36 / 60) train acc: 0.793000; val_acc: 0.601000
(Iteration 17801 / 29400) loss: 1.073276
(Iteration 18001 / 29400) loss: 0.973836
(Epoch 37 / 60) train acc: 0.811000; val_acc: 0.610000
(Iteration 18201 / 29400) loss: 0.973277
(Iteration 18401 / 29400) loss: 0.957541
(Iteration 18601 / 29400) loss: 0.991387
(Epoch 38 / 60) train acc: 0.827000; val_acc: 0.601000
(Iteration 18801 / 29400) loss: 0.745107
(Iteration 19001 / 29400) loss: 0.799122
(Epoch 39 / 60) train acc: 0.811000; val_acc: 0.598000
(Iteration 19201 / 29400) loss: 0.942136
(Iteration 19401 / 29400) loss: 0.800624

(Epoch 40 / 60) train acc: 0.783000; val_acc: 0.604000
(Iteration 19601 / 29400) loss: 0.940429
(Iteration 19801 / 29400) loss: 0.904573
(Iteration 20001 / 29400) loss: 0.993399
(Epoch 41 / 60) train acc: 0.816000; val_acc: 0.591000
(Iteration 20201 / 29400) loss: 0.899653
(Iteration 20401 / 29400) loss: 0.916943
(Epoch 42 / 60) train acc: 0.806000; val_acc: 0.596000
(Iteration 20601 / 29400) loss: 0.758303
(Iteration 20801 / 29400) loss: 0.778680
(Iteration 21001 / 29400) loss: 0.947486
(Epoch 43 / 60) train acc: 0.811000; val_acc: 0.592000
(Iteration 21201 / 29400) loss: 0.712344
(Iteration 21401 / 29400) loss: 0.806183
(Epoch 44 / 60) train acc: 0.826000; val_acc: 0.605000
(Iteration 21601 / 29400) loss: 0.783382
(Iteration 21801 / 29400) loss: 0.891782
(Iteration 22001 / 29400) loss: 0.834167
(Epoch 45 / 60) train acc: 0.819000; val_acc: 0.606000
(Iteration 22201 / 29400) loss: 0.797323
(Iteration 22401 / 29400) loss: 0.844305
```

```
(Epoch 46 / 60) train acc: 0.854000; val_acc: 0.614000
(Iteration 22601 / 29400) loss: 0.890112
(Iteration 22801 / 29400) loss: 0.944924
(Iteration 23001 / 29400) loss: 0.606655
(Epoch 47 / 60) train acc: 0.836000; val_acc: 0.606000
(Iteration 23201 / 29400) loss: 0.640090
(Iteration 23401 / 29400) loss: 0.970375
(Epoch 48 / 60) train acc: 0.830000; val_acc: 0.594000
(Iteration 23601 / 29400) loss: 0.662647
(Iteration 23801 / 29400) loss: 0.762254
(Iteration 24001 / 29400) loss: 0.780637
(Epoch 49 / 60) train acc: 0.842000; val_acc: 0.611000
(Iteration 24201 / 29400) loss: 0.901651
(Iteration 24401 / 29400) loss: 0.732053
(Epoch 50 / 60) train acc: 0.835000; val_acc: 0.601000
(Iteration 24601 / 29400) loss: 0.806242
(Iteration 24801 / 29400) loss: 1.107829
(Epoch 51 / 60) train acc: 0.844000; val_acc: 0.602000
(Iteration 25001 / 29400) loss: 0.814229
(Iteration 25201 / 29400) loss: 0.809749
(Iteration 25401 / 29400) loss: 0.762215
(Epoch 52 / 60) train acc: 0.852000; val_acc: 0.601000
(Iteration 25601 / 29400) loss: 0.823286
(Iteration 25801 / 29400) loss: 0.756253
(Epoch 53 / 60) train acc: 0.838000; val_acc: 0.605000

(Iteration 26001 / 29400) loss: 1.023917
(Iteration 26201 / 29400) loss: 0.755779
(Iteration 26401 / 29400) loss: 0.876049
(Epoch 54 / 60) train acc: 0.833000; val_acc: 0.598000
(Iteration 26601 / 29400) loss: 0.739671
(Iteration 26801 / 29400) loss: 0.771005
(Epoch 55 / 60) train acc: 0.838000; val_acc: 0.600000
(Iteration 27001 / 29400) loss: 0.625780
(Iteration 27201 / 29400) loss: 0.734696
(Iteration 27401 / 29400) loss: 0.845899
(Epoch 56 / 60) train acc: 0.869000; val_acc: 0.603000
(Iteration 27601 / 29400) loss: 0.761696
(Iteration 27801 / 29400) loss: 0.638063
(Epoch 57 / 60) train acc: 0.851000; val_acc: 0.607000
(Iteration 28001 / 29400) loss: 0.811818
(Iteration 28201 / 29400) loss: 0.958917
(Iteration 28401 / 29400) loss: 1.009080
(Epoch 58 / 60) train acc: 0.869000; val_acc: 0.603000
(Iteration 28601 / 29400) loss: 0.863148
(Iteration 28801 / 29400) loss: 0.497598
(Epoch 59 / 60) train acc: 0.850000; val_acc: 0.604000
(Iteration 29001 / 29400) loss: 0.876637
(Iteration 29201 / 29400) loss: 0.776647
(Epoch 60 / 60) train acc: 0.842000; val_acc: 0.612000
```

In [10]:

```
y_val_pred = np.argmax(model.loss(data['X_val']), axis = 1)
y_test_pred = np.argmax(model.loss(data['X_test']), axis = 1)
print('Validation accuracy is:', np.mean(y_val_pred == data['y_val']))
print('Testing accuracy is:', np.mean(y_test_pred == data['y_test']))
```

```
Validation accuracy is: 0.604
Testing accuracy is: 0.612
```

In [ ]:

```

```

In [ ]:

```

```